

Monitors

Loïc Guégan

Adapted from J. Kubiatowicz © 2010 UCB, O. J. Anshus and T. Larsen and P. Ha © UiT,
A. S. Tanenbaum © 2008, A. Silberschatz © 2009)

UiT The Arctic University of Norway

Spring - 2024

Outline

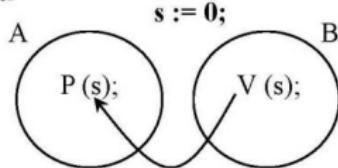
- **Motivation**
- Definition
- Monitor vs. semaphore
- Readers-Writers problem
 - ▶ A solution using monitors
- Language support for synchronization

Overview of Synchronization Mechanisms

Programs	Concurrent Programs				
Higher-level API	Shared Variables			Message Passing	
Hardware	Locks	Semaphores	Monitors	Send/Receive	
	Load/Store	Disable Ints	Test&Set	Comp&Swap	

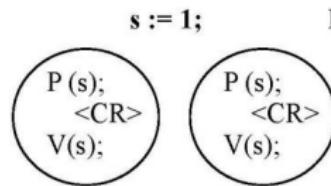
Review: semaphore

“The Signal”



A blocks until B says V

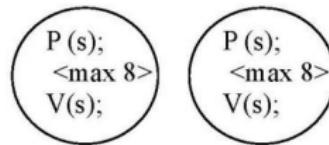
“The Mutex”



One thread gets in, next blocks until V is executed

NB: remember to set the initial semaphore value!

“The Team”



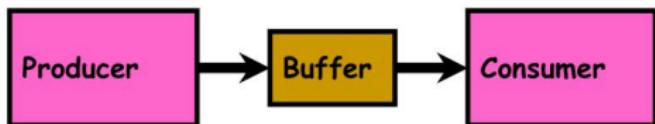
Up to 8 threads can pass P, the ninth will block until V is said by one of the eight already in there

Problems with Semaphores

Incorrect use of semaphore operations:

- $V(\text{mutex}) \dots P(\text{mutex})$
- $P(\text{mutex}) \dots V(\text{mutex})$
- Omitting of $P(\text{mutex})$ or $V(\text{mutex})$ (or both)

Recall: Producer-consumer with a bounded buffer



- Problem Definition
 - Producer puts things into a shared buffer
 - Consumer takes them out
 - Need synchronization to coordinate producer/consumer
- Don't want producer and consumer to have to work in lockstep, so put a fixed-size buffer between them
 - Need to synchronize access to this buffer
 - Producer needs to wait if buffer is full
 - Consumer needs to wait if buffer is empty
- Example: Coke machine
 - Producer can put limited number of cokes in machine
 - Consumer can't take cokes out if machine is empty



Recall: Solution to Bounded Buffer

```
Semaphore fullBuffer = 0;      // Initially, no item
Semaphore emptyBuffers = numBuffers;
                                // Initially, num empty slots
Semaphore mutex = 1;          // No one using machine

Producer(item) {
    emptyBuffers.P();           // Wait until space
    mutex.P();                  // Wait until buffer free
    Enqueue(item);
    mutex.V();
    fullBuffers.V();            // Tell consumers there is
                                // more item
}
Consumer() {
    fullBuffers.P();            // Check if there's an item
    mutex.P();                  // Wait until machine free
    item = Dequeue();
    mutex.V();
    emptyBuffers.V();           // tell producer need more
    return item;
}
```

Recall: Is the order of P's important?

```
Semaphore fullBuffer = 0;      // Initially, no item
Semaphore emptyBuffers = numBuffers;
                                         // Initially, num empty slots
Semaphore mutex = 1;           // No one using machine

Producer(item) {
    mutex.P();                  // Wait until buffer free
    emptyBuffers.P();           // Wait until space
    Enqueue(item);
    mutex.V();
    fullBuffers.V();           // Tell consumers there is
                               // more item
}
Consumer() {
    fullBuffers.P();            // Check if there's an item
    mutex.P();                  // Wait until machine free
    item = Dequeue();
    mutex.V();
    emptyBuffers.V();           // tell producer need more
    return item;
}
```

Motivation for Monitors and Condition Variables

- Semaphores are a huge step up, but:
 - They are confusing because they are dual purpose:
 - Both mutual exclusion and scheduling constraints
 - Example: the fact that flipping of P's in bounded buffer gives deadlock is not immediately obvious
 - Cleaner idea: Use **locks** for mutual exclusion and **condition variables** for scheduling constraints
- Definition: **Monitor**: a lock and zero or more condition variables for managing concurrent access to shared data
 - Use of Monitors is a programming paradigm
 - Some languages like Java provide monitors in the language

Outline

- Motivation
- **Definition**
- Monitor vs. semaphore
- Readers-Writers problem
 - ▶ A solution using monitors
- Language support for synchronization

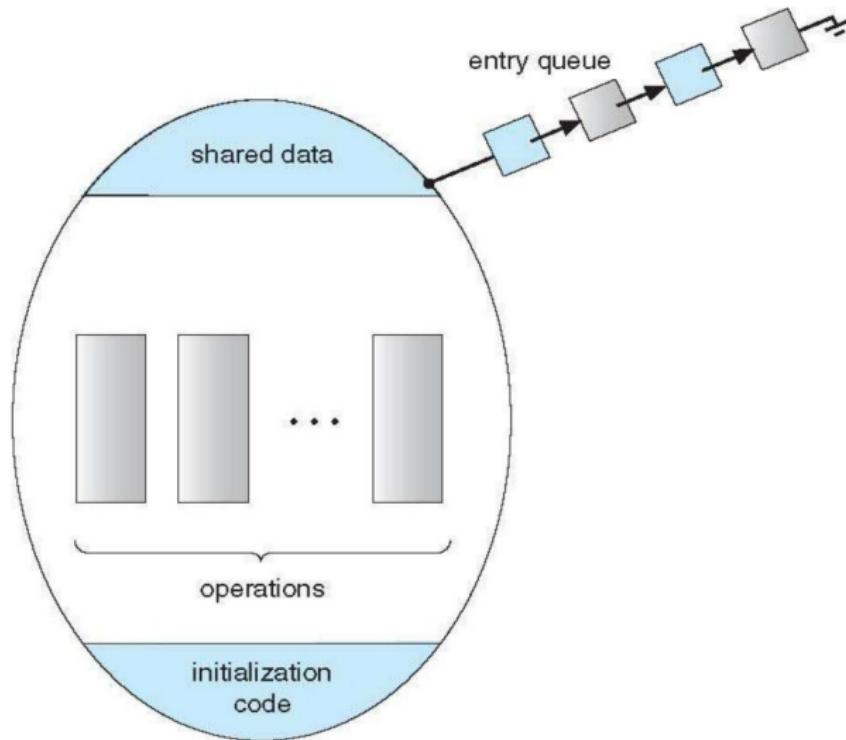
Monitors

Only one process may be active within the monitor at a time

```
monitor monitor-name
{
    // shared variable declarations
    procedure P1 (...) { .... }
    ...
    procedure Pn (...) {.....}
    Initialization code ( .... ) { ... }

}
```

Schematic view of a Monitor



Simple Monitor Example (version 1)

- Producer-consumer with an **infinite** buffer

```
Lock lock;
Queue queue;

AddToQueue(item) {           // Producer
    lock.Acquire();          // Lock shared data
    queue.enqueue(item);     // Add item
    lock.Release();          // Release Lock
}

RemoveFromQueue() {          // Consumer
    lock.Acquire();          // Lock shared data
    item = queue.dequeue();  // Get next item or null
    lock.Release();          // Release Lock
    return(item);            // Might return null
}
```

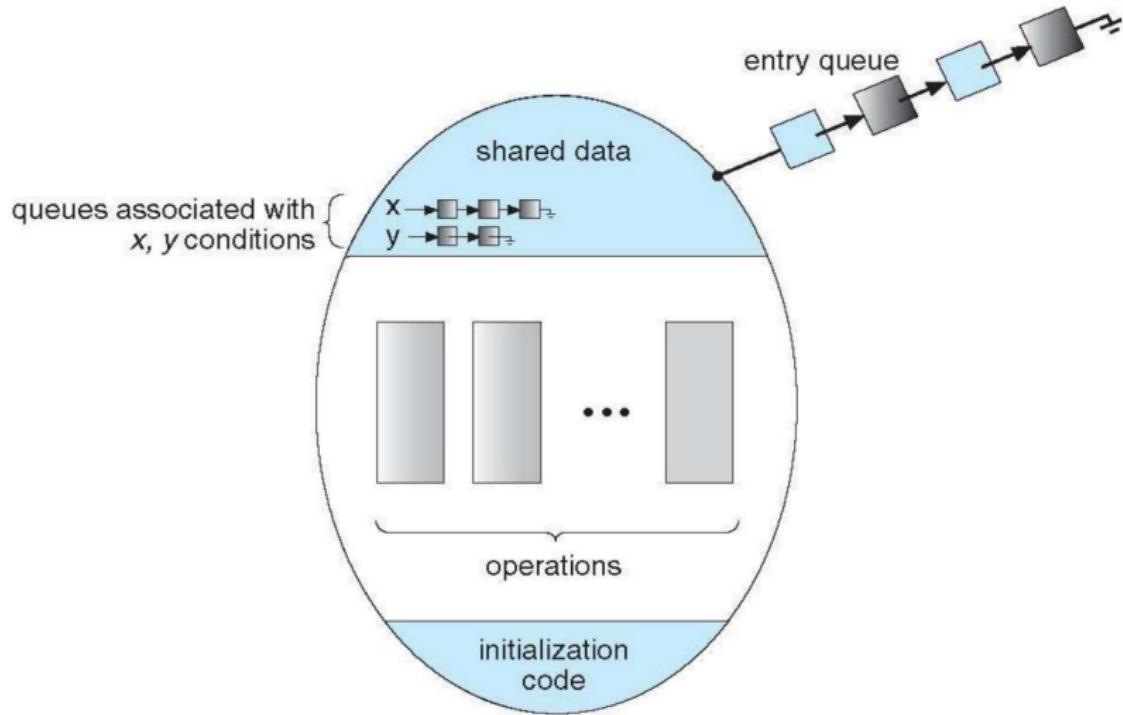
- Cannot put consumer to sleep if no work!

- Not very interesting use of "Monitor"
- It only uses a lock with no condition variables

Condition Variables

- How do we change the RemoveFromQueue() routine to wait until something is on the queue?
 - Could do this by keeping a count of the number of things on the queue (with semaphores), but error prone
- Condition Variable: a queue of threads waiting for something *inside* a critical section
 - Key idea: allow sleeping inside critical section by **atomically** releasing lock at time we go to sleep
 - Contrast to semaphores: Can't wait inside critical section
- Operations:
 - **Wait (&lock)**: Atomically release lock and go to sleep. Re-acquire lock later, before returning.
 - **Signal ()**: Wake up one waiter, if any
 - **Broadcast ()**: Wake up all waiters
- Rule: Must hold lock when doing condition variable ops!

Monitor with Condition Variables



Complete Monitor Example (with condition variable)

- Here is an (infinite) synchronized queue

```
Lock lock;
Condition dataready;
Queue queue;

AddToQueue(item) {
    lock.Acquire();                      // Get Lock
    queue.enqueue(item);                  // Add item
    dataready.signal();                  // Signal any waiters
    lock.Release();                     // Release Lock
}

RemoveFromQueue() {
    lock.Acquire();                      // Get Lock
    while (queue.isEmpty()) {
        dataready.wait(&lock);          // If nothing, sleep
    }
    item = queue.dequeue();              // Get next item
    lock.Release();                     // Release Lock
    return(item);
}
```

Mesa vs. Hoare monitors

- Need to be careful about precise definition of signal and wait. Consider a piece of our dequeue code:

```
while (queue.isEmpty()) {  
    dataready.wait(&lock); // If nothing, sleep  
}  
item = queue.dequeue(); // Get next item
```

- Why didn't we do this?

```
if (queue.isEmpty()) {  
    dataready.wait(&lock); // If nothing, sleep  
}  
item = queue.dequeue(); // Get next item
```

- Answer: depends on the type of scheduling

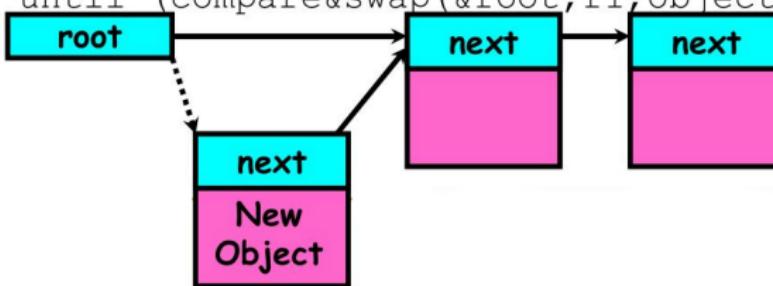
- Hoare-style (most textbooks):
 - Signaler gives lock, CPU to waiter; waiter runs immediately
 - Waiter gives up lock, processor back to signaller when it exits critical section or if it waits again
 - Mesa-style (most real operating systems):
 - Signaler keeps lock and processor
 - Waiter placed on ready queue with no special priority
 - Practically, need to check condition again after wait

Using of Compare & Swap for queues

- compare&swap (&address, reg1, reg2)
atomic {
 if (reg1 == M[address]) {
 M[address] = reg2;
 return success;
 } else {
 return failure;
 }
}

Here is an atomic add to linked-list function:

```
addToQueue (&object) {  
    do {  
        ld r1, M[root]      // repeat until no conflict  
        st r1, M[object]    // Get ptr to current head  
        st r1, M[object]    // Save link in new object  
    } until (compare&swap(&root, r1, object));  
}
```



Outline

- Motivation
- Definition
- **Monitor vs. semaphore**
- Readers-Writers problem
 - ▶ A solution using monitors
- Language support for synchronization

Can we construct Monitors from Semaphores?

- Locking aspect is easy: Just use a mutex
- Can we implement condition variables this way?

```
Wait() { semaphore.P(); }
Signal() { semaphore.V(); }
```

- Does this work better?

```
Wait(Lock lock) {
    lock.Release();
    semaphore.P();
    lock.Acquire();
}
Signal() { semaphore.V(); }
```

Construction of Monitors from Semaphores (con't)

- Problem with previous try:
 - P and V are commutative - result is the same no matter what order they occur
 - Condition variables are NOT commutative
- Does this fix the problem?

```
Wait(Lock lock) {  
    lock.Release();  
    semaphore.P();  
    lock.Acquire();  
}  
Signal() {  
    if semaphore queue is not empty  
        semaphore.V();  
}
```

- Not legal to look at contents of semaphore queue
- There is a race condition - signaler can slip in after lock release and before waiter executes semaphore.P()
- It is actually possible to do this correctly
 - Complex solution for Hoare monitor in book Operating System Concepts by A. Silberschatz.

Semaphore vs. Monitor

Semaphore

- **P(s)** means WAIT if $s = 0$.
Otherwise, $s = s - 1$
- **V(s)** means start a waiting thread and REMEMBER that a V call was made: $s = s + 1$
- Assume $s = 0$ when V(s) is called: If there is no thread to start this time, the next thread to call P(s) will get through P(s)

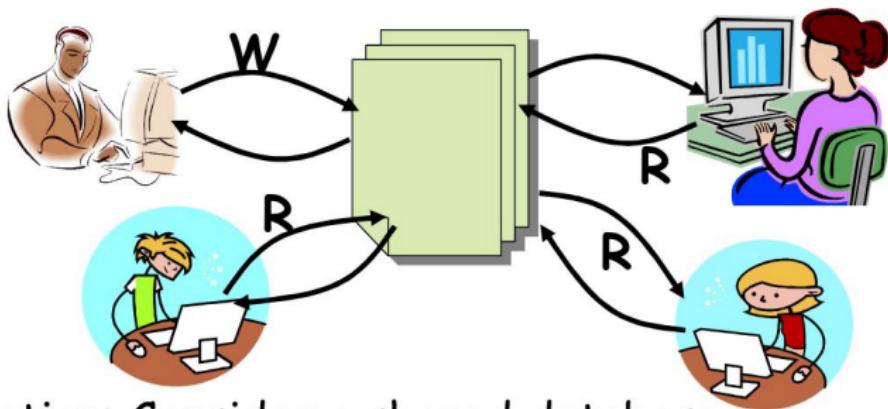
Monitor

- **Wait(cond)** means unconditional WAIT
- **Signal(cond)** means start a waiting thread. But no history!
- Assume that the condition queue is empty when Signal() is called. The next thread to call Wait(cond) will block.

Outline

- Motivation
- Definition
- Monitor vs. semaphore
- **Readers-Writers problem**
 - ▶ A solution using monitors
- Language support for synchronization

Readers/Writers Problem



- Motivation: Consider a shared database
 - Two classes of users:
 - Readers - never modify database
 - Writers - read and modify database
 - Is using a single lock on the whole database sufficient?
 - Like to have many readers at the same time
 - Only one writer at a time

Basic Readers/Writers Solution

■ Correctness Constraints:

- Readers can access database when no writers
- Writers can access database when no readers or writers
- Only one thread manipulates state variables at a time

■ Basic structure of a solution:

- **Reader()**

- Wait until no writers

- Access data base

- Check out - wake up a waiting writer

- **Writer()**

- Wait until no **active** readers or writers

- Access database

- Check out - wake up waiting readers or writer

- State variables (Protected by a lock called "lock"):

- int AR: Number of active readers; initially = 0

- int WR: Number of waiting readers; initially = 0

- int AW: Number of active writers; initially = 0

- int WW: Number of waiting writers; initially = 0

- Condition okToRead = NIL

- Condition okToWrite = NIL

Code for a Reader

```
Reader() {
    // First check self into system
    lock.Acquire();
    while ((AW + WW) > 0) {      // Is it safe to read?
        WR++; // No. Writers exist
        okToRead.wait(&lock);    // Sleep on cond var
        WR--; // No longer waiting
    }
    AR++;                      // Now we are active!
    lock.release();
    // Perform actual read-only access
    AccessDatabase(ReadOnly);
    // Now, check out of system
    lock.Acquire();
    AR--;                      // No longer active
    if (AR == 0 && WW > 0)      // No other active readers
        okToWrite.signal();    // Wake up one writer
    lock.Release();
}
```

Code for a Writer

```
Writer() {
    // First check self into system
    lock.Acquire();
    while ((AW + AR) > 0) { // Is it safe to write?
        WW++; // No. Active users exist
        okToWrite.wait(&lock); // Sleep on cond var
        WW--; // No longer waiting
    }
    AW++; // Now we are active!
    lock.release();
    // Perform actual read/write access
    AccessDatabase(ReadWrite);
    // Now, check out of system
    lock.Acquire();
    AW--; // No longer active
    if (WW > 0){ // Give priority to writers
        okToWrite.signal(); // Wake up one writer
    } else if (WR > 0) { // Otherwise, wake reader
        okToRead.broadcast(); // Wake all readers
    }
    lock.Release();
}
```

Simulation of Readers/Writers solution

- Consider the following sequence of operators:
 - R1, R2, W1, R3
- On entry, each reader checks the following:

```
while ((AW + WW) > 0) {           // Is it safe to read?  
    WR++;                         // No. Writers exist  
    okToRead.wait(&lock);          // Sleep on cond var  
    WR--;                          // No longer waiting  
}  
AR++;                           // Now we are active!
```

- First, R1 comes along:
 $AR = 1, WR = 0, AW = 0, WW = 0$
- Next, R2 comes along:
 $AR = 2, WR = 0, AW = 0, WW = 0$
- Now, readers may take a while to access database
 - Situation: Locks released

Simulation(2)

- Next, W1 comes along:

```
while ((AW + AR) > 0) { // Is it safe to write?  
    WW++; // No. Active users exist  
    okToWrite.wait(&lock); // Sleep on cond var  
    WW--; // No longer waiting  
}  
AW++;
```

- Can't start because of readers, so go to sleep:

AR = 2, WR = 0, AW = 0, WW = 1

- Finally, R3 comes along:

AR = 2, WR = 1, AW = 0, WW = 1

- Now, say that R2 finishes before R1:

AR = 1, WR = 1, AW = 0, WW = 1

- Finally, last of first two readers (R1) finishes and wakes up writer:

```
if (AR == 0 && WW > 0) // No other active readers  
    okToWrite.signal(); // Wake up one writer
```

Simulation(3)

- When writer wakes up, get:

$$AR = 0, WR = 1, AW = 1, WW = 0$$

- Then, when writer finishes:

```
if (WW > 0) {           // Give priority to writers
    okToWrite.signal();   // Wake up one writer
} else if (WR > 0) {      // Otherwise, wake reader
    okToRead.broadcast(); // Wake all readers
}
```

- Writer wakes up reader, so get:

$$AR = 1, WR = 0, AW = 0, WW = 0$$

- When reader completes, we are finished

Questions

- Can readers starve? Consider Reader() entry code:

```
while ((AW + WW) > 0) {      // Is it safe to read?  
    WR++;                      // No. Writers exist  
    okToRead.wait(&lock);       // Sleep on cond var  
    WR--;                      // No longer waiting  
}  
  
AR++;                         // Now we are active!
```

- What if we erase the condition check in Reader exit?

```
AR--;                          // No longer active  
if (AR == 0 && WW > 0)          // No other active readers  
    okToWrite.signal();         // Wake up one writer
```

- Further, what if we turn the signal() into broadcast()

```
AR--;                          // No longer active  
okToWrite.broadcast();         // Wake up one writer
```

- Finally, what if we use only one condition variable (call it "okToContinue") instead of two separate ones?

- ❑ Both readers and writers sleep on this variable
 - ❑ Must use broadcast() instead of signal()

Monitor Conclusion

- Monitors represent the logic of the program
 - Wait if necessary
 - Signal when change something so any waiting threads can proceed
- Basic structure of monitor-based program:

```
lock  
while (need to wait) {  
    condvar.wait();  
}  
unlock
```

Check and/or update state variables

Wait if necessary

do something so no need to wait

```
lock  
condvar.signal();  
unlock
```

Check and/or update state variables

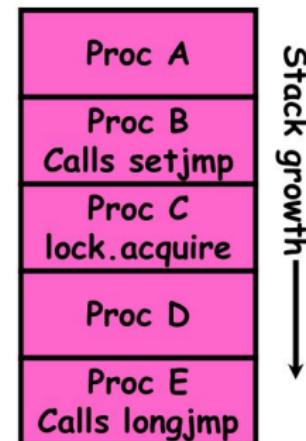
Outline

- Motivation
- Definition
- Monitor vs. semaphore
- Readers-Writers problem
 - ▶ A solution using monitors
- **Language support for synchronization**

C-Language Support for Synchronization

- C language: Pretty straightforward synchronization
 - Just make sure you know *all* the code paths out of a critical section

```
int Rtn() {  
    lock.acquire();  
    ...  
    if (exception) {  
        lock.release();  
        return errReturnCode;  
    }  
    ...  
    lock.release();  
    return OK;  
}
```



- Watch out for `setjmp/longjmp`
 - Can cause a non-local jump out of procedure
 - In example, procedure E calls `longjmp`, popping stack back to procedure B
 - If Procedure C had `lock.acquire`, problem!

C++ Support for Synchronization

- Languages with exceptions like C++
 - Languages that support exceptions are problematic (easy to make a non-local exit without releasing lock)
 - Consider:

```
void Rtn() {  
    lock.acquire();  
    ...  
    DoFoo();  
    ...  
    lock.release();  
}  
void DoFoo() {  
    ...  
    if (exception) throw errException;  
    ...  
}
```

- Notice that an exception in DoFoo() will exit without releasing the lock
-

C++ Support for Synchronization (con't)

- Must catch all exceptions in critical sections
 - Catch exceptions, release lock, and re-throw exception:

```
void Rtn() {  
    lock.acquire();  
    try {  
        ...  
        DoFoo();  
        ...  
    } catch (...) {          // catch exception  
        lock.release();      // release lock  
        throw;                // re-throw the exception  
    }  
    lock.release();  
}  
void DoFoo() {  
    ...  
    if (exception) throw errException;  
    ...  
}
```

References

- A. S. Tanenbaum, Modern Operating Systems.
- A. Silberschatz et. al., Operating System Concepts.

Thanks for your attention!

Questions?