

Deadlocks

Loïc Guégan

Adapted from J. Kubiawicz @ 2010 UCB, O. J. Anshus and T. Larsen and P. Ha @ UiT,
K. Li @ Princeton, A. S. Tanenbaum @ 2008, A. Silberschatz @ 2009

UiT The Arctic University of Norway

Spring - 2024

Outline

- What is deadlock?
- How can deadlock occur?
- How to deal with deadlocks?
- Examples




Resources

- Resources - passive entities needed by threads to do their work
 - ▶ CPU time, disk space, memory
- Two types of resources:
 - 1 **Preemptable** - can take away
 - ★ e.g. CPU
 - 2 **Non-preemptable** - must leave it with the thread
 - ★ Disk space, chunk of virtual address space
 - ★ Mutual exclusion - the right to enter a critical section
- Resources may require exclusive access or may be sharable
 - ▶ Read-only files are typically sharable
 - ▶ Printers are not sharable during time of printing
- One of the major tasks of an operating system is to **manage resources**



Definition: Starvation vs Deadlock


- Starvation \Rightarrow thread waits indefinitely
 - ▶ Example, low-priority thread waiting for resources constantly in use by high-priority threads
- Deadlock \Rightarrow circular waiting for resources
 - ▶ Thread A owns Res 1 and is waiting for Res 2
 - ▶ Thread B owns Res 2 and is waiting for Res 1



- Deadlock vs Starvation
 - ▶ Starvation can end (but do not have to)
 - ▶ Deadlock cannot end without external intervention

An example

- A utility program
 - Copy a file from tape to disk
 - Print the file to printer
- Resources
 - Tape
 - Disk
 - Printer
- A deadlock
 - A holds tape and disk, then requests printer
 - B holds printer, then requests tape and disk



Bridge crossing example



- Each segment of road can be viewed as a resource
 - Cars must own the segment under them
 - Must acquire segment that they are moving into
- For bridge: must acquire both halves
 - Traffic only in one direction at a time
 - Problem occurs when two cars in opposite directions on bridge: each acquires one segment and needs next
- If a deadlock occurs, it can be resolved if one car backs up (release resources and rollback)
 - Several cars may have to be backed up
- Starvation is possible
 - East-going traffic really fast \Rightarrow no one goes west

Deadlock properties

- Deadlock **not always deterministic** - Example 2 mutexes:

Thread A

`x.P();`

`y.P();`

`y.V();`

`x.V();`

Thread B

`y.P();`

`x.P();`

`x.V();`

`y.V();`

- Deadlock won't always happen with this code
 - Have to have exactly the right timing ("wrong" timing?)
- Deadlocks **occur with multiple resources**
 - Can't decompose the problem
 - Can't solve deadlock for each resource independently
- Example: System with 2 disk drives and two threads
 - Each thread needs 2 disk drives to function
 - Each thread gets one disk and waits for another one

Resource Allocation Graph


- Process A is holding resource R



- Process B is wait for resource S




- A cycle in a resource allocation graph \rightarrow deadlock
- If A waits for S while holding R, and B waits for R while holding S, then




Deadlock Modeling

1. A requests R
2. B requests S
3. C requests T
4. A requests S
5. B requests T
6. C requests R
deadlock


(d)




(e)




(f)




(g)



(h)



(i)




(j)


Deadlock Modeling

1. A requests R
 2. C requests T
 3. A requests S
 4. C requests R
 5. A releases R
 6. A releases S
- no deadlock


(k)




(l)




(m)




(n)



(o)



(p)




(q)

Multiple instance of a resource

■ System Model

- A set of Threads T_1, T_2, \dots, T_n
- Resource types R_1, R_2, \dots, R_m
CPU cycles, memory space, I/O devices
- Each resource type R_i has **W_i instances**.
- Each thread utilizes a resource as follows:
 - Request() / Use() / Release()




■ Resource-Allocation Graph:

- V is partitioned into two types:
 - $T = \{T_1, T_2, \dots, T_n\}$, the set threads in the system.
 - $R = \{R_1, R_2, \dots, R_m\}$, the set of resource types in system
- Request edge: directed edge $T_i \rightarrow R_j$
- Assignment edge: directed edge $R_j \rightarrow T_i$


Resource Allocation Graph examples

■ Recall:


- request edge - directed edge $T_i \rightarrow R_j$
- assignment edge - directed edge $R_j \rightarrow T_i$



Simple Resource
Allocation Graph




Allocation Graph
With Deadlock



Allocation Graph
With Cycle, but
No Deadlock


Necessary conditions for deadlock

- Mutual exclusion condition
- Hold and wait
- No resource preemption
- Circular waiting




Eliminate competition for resources?

- If running A to completion and then running B, there will be no deadlock.
- Should CPU scheduling algorithm eliminate competition for resources?



First iteration




Second iteration

- What is deadlock?
- How can deadlock occur?
- **Strategies to deal with deadlocks?**
 - ▶ Ignore the problem
 - ▶ Detection and recovery
 - ▶ Dynamic avoidance
 - ▶ Prevention
- Examples


Strategies

- Ignore the problem
 - ▶ It is user's fault
- Detection and recovery
 - ▶ Fix the problem afterwards
- Dynamic avoidance
 - ▶ Careful allocation
- Prevention
 - ▶ Remove one of the four conditions

Deadlock detection with ONE resource of each type (1)



(a)



(b)

(a) A resource graph. (b) A cycle extracted from (a).

Deadlock detection algorithm

- Only one of each type of resource \Rightarrow look for loops
- More General Deadlock Detection Algorithm


- Let $[X]$ represent an m -ary vector of non-negative integers (quantities of resources of each type), e.g. $[2,2]$:

$[FreeResources]$: Current free resources each type, e.g. $[0,0]$
 $[Request_x]$: Current requests from thread X , e.g. $[1,0]$ for T_1
 $[Alloc_x]$: Current resources held by thread X , $[0,1]$ for T_1

- See if tasks can eventually terminate on their own

```
[Avail] = [FreeResources]
Add all nodes to UNFINISHED
do {
    done = true
    Foreach node in UNFINISHED {
        if ([Requestnode] <= [Avail]) {
            remove node from UNFINISHED
            [Avail] = [Avail] + [Allocnode]
            done = false
        }
    }
} until(done)
```

- Nodes left in UNFINISHED \Rightarrow deadlocked



Deadlock detection with MULTIPLE resource of each type

<div><div><div><i>Tape drives</i></div><div><i>Plotters</i></div><div><i>Scanners</i></div><div><i>CD Roms</i></div></div><div>E = (4 2 3 1)</div><div>Resources in existence</div></div>	<div><div><div><i>Tape drives</i></div><div><i>Plotters</i></div><div><i>Scanners</i></div><div><i>CD Roms</i></div></div><div>A = (2 1 0 0)</div><div>Resources available</div></div>
<div><div><div>C =</div><div><div><div>0010</div><div>2001</div><div>0120</div></div><div><div>Thread 1</div><div>Thread 2</div><div>Thread 3</div></div></div></div><div>Current allocation matrix</div></div>	<div><div><div>R =</div><div><div><div>2001</div><div>1010</div><div>2100</div></div><div><div>Thread 1</div><div>Thread 2</div><div>Thread 3</div></div></div></div><div>Request matrix</div></div>

Deadlock detection with MULTIPLE resource of each type

<i>Tape drives</i>	<i>Plotters</i>	<i>Scanners</i>	<i>CD Roms</i>		<i>Tape drives</i>	<i>Plotters</i>	<i>Scanners</i>	<i>CD Roms</i>	
E = (4 2 3 1)				Resources in existence	A = (2 2 2 0)				Resources available
C = $\begin{bmatrix} 0 & 0 & 1 & 0 \\ 2 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 \end{bmatrix}$					R = $\begin{bmatrix} 2 & 0 & 0 & 1 \\ 1 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 \end{bmatrix}$				
Thread 1					Thread 1				
Thread 2					Thread 2				
Thread 3					Thread 3				
Current allocation matrix					Request matrix				

Deadlock detection with MULTIPLE resource of each type

<div><div><div>Tape drives</div><div>Plotters</div><div>Scanners</div><div>CD Roms</div></div><div>E = (4 2 3 1)</div><div>Resources in existence</div></div>	<div><div><div>Tape drives</div><div>Plotters</div><div>Scanners</div><div>CD Roms</div></div><div>A = (4 2 2 1)</div><div>Resources available</div></div>
<div><div><div><div>C =</div><div><div><div>0</div><div>0</div><div>1</div><div>0</div></div><div><div>0</div><div>0</div><div>0</div><div>0</div></div><div><div>0</div><div>0</div><div>0</div><div>0</div></div></div></div><div><div>Thread 1</div><div>Thread 2</div><div>Thread 3</div></div></div><div>Current allocation matrix</div></div>	<div><div><div><div>R =</div><div><div><div>2</div><div>0</div><div>0</div><div>1</div></div><div><div>0</div><div>0</div><div>0</div><div>0</div></div><div><div>0</div><div>0</div><div>0</div><div>0</div></div></div></div><div><div>Thread 1</div><div>Thread 2</div><div>Thread 3</div></div></div><div>Request matrix</div></div>

Deadlock detection with MULTIPLE resource of each type

<div><div><div>Tape drives</div><div>Plotters</div><div>Scanners</div><div>CD Roms</div></div><div>E = (4 2 3 1)</div><div>Resources in existence</div></div>	<div><div><div>Tape drives</div><div>Plotters</div><div>Scanners</div><div>CD Roms</div></div><div>A = (4 2 3 1)</div><div>Resources available</div></div>
<div><div><div>C =</div><div><div><div>0</div><div>0</div><div>0</div><div>0</div></div><div>Thread 1</div></div><div><div><div>0</div><div>0</div><div>0</div><div>0</div></div><div>Thread 2</div></div><div><div><div>0</div><div>0</div><div>0</div><div>0</div></div><div>Thread 3</div></div></div></div> <div>Current allocation matrix</div>	<div><div><div>R =</div><div><div><div>0</div><div>0</div><div>0</div><div>0</div></div><div>Thread 1</div></div><div><div><div>0</div><div>0</div><div>0</div><div>0</div></div><div>Thread 2</div></div><div><div><div>0</div><div>0</div><div>0</div><div>0</div></div><div>Thread 3</div></div></div></div> <div>Request matrix</div>

- Kill process/thread
 - ▶ Can you always do this?
- Preempt resources without killing threads
 - ▶ What if the resources are in a critical section?
- Roll back actions of deadlocked threads
 - ▶ Like transactions in databases

- What is deadlock?
- How can deadlock occur?
- **Strategies to deal with deadlocks?**
 - ▶ Ignore the problem
 - ▶ Detection and recovery
 - ▶ **Dynamic avoidance**
 - ▶ Prevention
- Examples

- Model
 - ▶ Each thread requests resources **one at a time**
- Safety conditions:
 - ▶ It is not deadlocked
 - ★ Found scheduling order in which every thread can run to completion (even if all request their max resources)

Examples (single resource)

Total: 8

	Has	Max
P ₁	2	6
P ₂	2	3
P ₃	3	5

Free: 1

	Has	Max
P ₁	2	6
P ₂	3	3
P ₃	3	5

Free: 0

	Has	Max
P ₁	2	6
P ₂	0	0
P ₃	3	5

Free: 3

	Has	Max
P ₁	2	6
P ₂	0	0
P ₃	5	5

Free: 1

	Has	Max
P ₁	2	6
P ₂	0	0
P ₃	0	0

Free: 6

	Has	Max
P ₁	4	6
P ₂	1	3
P ₃	2	5

?

Free: 1

Banker's algorithm (multiple resources)

■ Toward right idea:

- State maximum resource needs in advance
- Allow particular thread to proceed if:
(available resources - #requested) \geq max
remaining that might be needed by **any** thread

■ Banker's algorithm (less conservative):

- Allocate resources dynamically
 - Evaluate each request and grant if threads are still deadlock free afterward
 - Technique: pretend each request is granted, then run deadlock detection algorithm, substituting $([Max_{node}] - [Alloc_{node}] \leq [Avail])$ for $([Request_{node}] \leq [Avail])$
Grant request if result is deadlock free (conservative!)
 - Keeps system in a "SAFE" state, i.e. there exists a sequence $\{T_1, T_2, \dots, T_n\}$ with T_1 requesting all remaining resources, finishing, then T_2 requesting all remaining resources, etc..



Safe state in the Banker's algorithm (single resources)

- Three resource allocation states: (a) Safe. (b) Safe. (c) Unsafe.

Has Max		
A	0	6
B	0	5
C	0	4
D	0	7

Free: 10

(a)



Has Max		
A	1	6
B	1	5
C	2	4
D	4	7

Free: 2

(b)



Has Max		
A	1	6
B	2	5
C	2	4
D	4	7


Free: 1

(c)

- What is deadlock?
- How can deadlock occur?
- **Strategies to deal with deadlocks?**
 - ▶ Ignore the problem
 - ▶ Detection and recovery
 - ▶ Dynamic avoidance
 - ▶ **Prevention**
- Examples

Prevention: avoid mutual exclusion

- Some can be made sharable
 - Read-only files, memory, etc
- Some resources are not physically sharable
 - Printer, tape, etc
 - Some can be virtualized by spooling
- What about the tape-disk-printer example?



Prevention: avoid hold and wait

■ Two-phase locking

□ Assumption

- Processes know all resources they will need at the beginning

□ Phase I:

- Try to lock all resources at the beginning

□ Phase II:


- If successful, use the resources and release them

- Otherwise, release all resources and start over

■ Application

- Telephone company's circuit switching


■ What about the tape-disk-printer example?



Prevention: allow preemption


- Make the scheduler be aware of resource allocation
 - Method
 - If the system cannot satisfy a request from a thread holding resources, preempt the thread and release its resources
 - Schedule a thread only if the system satisfies all resources

- What about the tape-disk-printer example?




Prevention: avoid circular wait

- Impose an order of requests for all resources
- Method
 - Assign a unique id to each resource
 - All requests must be in ascending order of the ids
- What about the tape-disk-printer example?
- Does this method have no circular wait?



- What is deadlock?
- How can deadlock occur?
- Strategies to deal with deadlocks?
 - ▶ Ignore the problem
 - ▶ Detection and recovery
 - ▶ Dynamic avoidance
 - ▶ Prevention
- **Examples**
 - ▶ **Dining philosophers**

“Dining Philosophers”



s
s(i): One
semaphore per fork
to be used in
mutex style P-V

- Each: need 2 forks to eat
- 5 philosophers: 10 forks
- 5 forks: 2 can eat concurrently

Things to observe:

- A fork can only be used by one at a time, please
- No deadlock, please
- No starving, please
- Concurrent eating, please


T_i

T_i

T_i

Dining philosophers 1

“Dining Philosophers”



s
 s(i): One
 semaphore per fork
 to be used in
mutex style P-V

- Each: need 2 forks to eat
- 5 philosophers: 10 forks
- 5 forks: 2 can eat concurrently

Mutex on whole table: $P(mutex);$ T_i
 • *I can eat at a time* $eat;$
 $V(mutex);$

Things to observe:


- A fork can only be used by one at a time, please
- No deadlock, please
- No starving, please
- Concurrent eating, please

T_i

T_i

Dining philosophers 1

“Dining Philosophers”



S
s(i): One
semaphore per fork
to be used in
mutex style P-V

- Each: need 2 forks to eat
- 5 philosophers: 10 forks
- 5 forks: 2 can eat concurrently

Mutex on whole table:	$P(\text{mutex});$	T_i
• <i>I can eat at a time</i>	eat;	
	$V(\text{mutex});$	

Get L; Get R;	$P(s(i));$	T_i
• <i>Deadlock possible</i>	$P(s(i+1));$	
	eat;	
	$V(s(i+1));$	
S(i) = 1 initially	$V(s(i));$	


Things to observe:

- A fork can only be used by one at a time, please
- No deadlock, please
- No starving, please
- Concurrent eating, please

T_i

Dining philosophers 1

“Dining Philosophers”



s
s(i): One semaphore per fork to be used in mutex style P-V

- Each: need 2 forks to eat
- 5 philosophers: 10 forks
- 5 forks: 2 can eat concurrently

Things to observe:

- A fork can only be used by one at a time, please
- No deadlock, please
- No starving, please
- Concurrent eating, please

Mutex on whole table: $P(mutex);$ T_i
 $eat;$
 $V(mutex);$
•*I can eat at a time*


Get L; Get R; $P(s(i));$ T_i
 $P(s(i+1));$
 $eat;$
 $V(s(i+1));$
S(i) = 1 initially $V(s(i));$
•*Deadlock possible*

Avoid hold-and-wait

Get L; Get R if free else Put L; T_i
•*Starvation possible*

Dining philosophers 1

Dining Philosophers



Can we in a simple way do better than this one?

Get L; Get R;

•Deadlock possible

```
P(s(i));  
P(s(i+1));  
eat;  
V(s(i+1));  
V(s(i));
```

Avoid circular wait



$S(i) = 1$ initially

$T_1, T_2, T_3, T_4:$

```
P(s(i));  
P(s(i+1));  
<eat>  
V(s(i+1));  
V(s(i));
```

- Remove the danger of circular waiting (deadlock)
- T_1 - T_4 : Get L; Get R;
- T_5 : Get R; Get L;

T_5

```
P(s(1));  
P(s(5));  
<eat>  
V(s(5));  
V(s(1));
```

•Non-symmetric solution. Still quite elegant

Summary: strategies to deal with deadlocks

- Ignore the problem
 - ▶ It is user's fault
- Detection and recovery
 - ▶ Fix the problem afterwards
- Dynamic avoidance
 - ▶ Careful allocation
- Prevention (Negate one of the four conditions)
 - ▶ Avoid the mutual exclusion
 - ▶ Avoid the hold and wait
 - ▶ Allow preemption
 - ▶ Avoid the circular wait

References

- A. S. Tanenbaum, Modern Operating Systems.
- A. Silberschatz et. al., Operating System Concepts.

Thanks for your attention!

Questions?