

CPU Scheduling

Loïc Guégan

Adapted from P. Ha @ UiT, J. Kubiawicz @ 2010 UCB,
K. Li and J.P. Singh @ Princeton, A. S. Tanenbaum @ 2008, A. Silberschatz @ 2009

UiT The Arctic University of Norway

Spring - 2024

Outline

- What is CPU-scheduling?
- Why is CPU-scheduling needed?
- How does CPU-scheduling work?

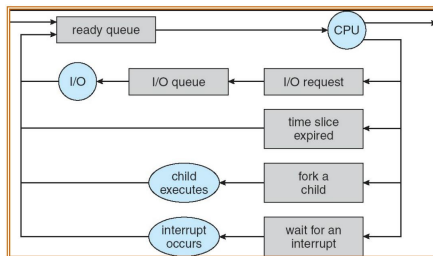
Recall: Dispatching Loop

- Conceptually, the dispatching loop of the operating system looks as follows:

```
Loop {  
    RunThread ();  
    ChooseNextThread ();  
    SaveStateOfCPU (curTCB );  
    LoadStateOfCPU (newTCB );  
}
```

- This is an infinite loop
 - ▶ One could argue that this is all that the OS does

CPU Scheduling



- How OS decides the threads to take off from the queues?
 - ▶ One is important \Rightarrow ready queue
 - ▶ Others can be scheduled as well but not the focus
- **CPU-scheduling:** deciding which threads are given access to CPU from moment to moment

Outline

- What is CPU-scheduling?
- Why is CPU-scheduling needed?
- How does CPU-scheduling work?
 - ▶ Scheduling in batch and interactive systems

Why CPU-scheduling?

- Multiple threads compete for CPU
- There aren't enough resources for all threads
- CPU time is a scarce resource

Scheduling Criteria

- Assumptions:
 - ▶ One program per user and one thread per program
 - ▶ Programs are independent
 - Goals for batch and interactive systems:
 - 1 Maximize throughput
 - ★ Jobs/ hour
 - ★ Min overhead, max resource utilization
 - 2 Minimize completion time
 - ★ Average time from submission to completion
 - 3 Minimize response time
 - ★ Time until response (e.g. typing on keyboard)
 - 4 Fairness
- } Batch systems
- } Interactive systems

- What is CPU-scheduling?
- Why is CPU-scheduling needed?
- How does CPU-scheduling work?
 - ▶ Scheduling in batch and interactive systems

First-Come First-Serve (FCFS) Policy

- What does it mean?
 - ▶ Serve in the order of requests
 - ▶ Run to completion or until blocked or yields
- Example 1
 - ▶ P1 = 24sec, P2 = 3sec, and P3 = 3sec, submitted together
 - ▶ Average completion time = $(24 + 27 + 30) / 3 = 27$




- Example 2
 - ▶ Same jobs but come in different order: P2, P3 and P1
 - ▶ Average completion time =




Scheduling – Process Behavior

- Bursts of CPU usage alternate with periods of waiting for I/O



(a) A CPU-bound process. (b) An I/O-bound process.

Round Robin




- Similar to FCFS, but add a time slice for timer interrupt
- FCFS for preemptive scheduling

Example of RR with Time Slice = 20s

- Example:

Thread	Burst Time (s)
P_1	53
P_2	8
P_3	68
P_4	24

- The Gantt chart is:



- Waiting time for

0 20 28 48 68 88 108 112 125 145 153

$$P_1 = (68 - 20) + (112 - 88) = 72$$

$$P_2 = (20 - 0) = 20$$

$$P_3 = (28 - 0) + (88 - 48) + (125 - 108) = 85$$

$$P_4 = (48 - 0) + (108 - 68) = 88$$


- Average waiting time = $(72 + 20 + 85 + 88)/4 = 66\frac{1}{4}$
- Average completion time = $(125 + 28 + 153 + 112)/4 = 104\frac{1}{2}$
- Thus, Round-Robin Pros and Cons:
 - ▶ Better for short jobs, Fair (+)
 - ▶ Context-switching time adds up for long jobs (-)

Example of RR with Time Slice = 20s

- Example:

Thread	Burst Time (s)
P_1	53
P_2	8
P_3	68
P_4	24

- The Gantt chart is:



- Waiting time for

0 20 28 48 68 88 108 112 125 145 153

$$P_1 = (68 - 20) + (112 - 88) = 72$$

$$P_2 = (20 - 0) = 20$$

$$P_3 = (28 - 0) + (88 - 48) + (125 - 108) = 85$$

$$P_4 = (48 - 0) + (108 - 68) = 88$$


- Average waiting time = $(72 + 20 + 85 + 88)/4 = 66\frac{1}{4}$
- Average completion time = $(125 + 28 + 153 + 112)/4 = 104\frac{1}{2}$
- Thus, Round-Robin Pros and Cons:
 - ▶ Better for short jobs, Fair (+)
 - ▶ Context-switching time adds up for long jobs (-)

Example of RR with Time Slice = 20s

- Example:

Thread	Burst Time (s)
P_1	53
P_2	8
P_3	68
P_4	24

- The Gantt chart is:



- Waiting time for

$$P_1 = (68 - 20) + (112 - 88) = 72$$

$$P_2 = (20 - 0) = 20$$

$$P_3 = (28 - 0) + (88 - 48) + (125 - 108) = 85$$

$$P_4 = (48 - 0) + (108 - 68) = 88$$


- Average waiting time = $(72 + 20 + 85 + 88)/4 = 66\frac{1}{4}$
- Average completion time = $(125 + 28 + 153 + 112)/4 = 104\frac{1}{2}$
- Thus, Round-Robin Pros and Cons:
 - ▶ Better for short jobs, Fair (+)
 - ▶ Context-switching time adds up for long jobs (-)

Example of RR with Time Slice = 20s

- Example:

Thread	Burst Time (s)
P_1	53
P_2	8
P_3	68
P_4	24

- The Gantt chart is:



- Waiting time for

0 20 28 48 68 88 108 112 125 145 153

$$P_1 = (68 - 20) + (112 - 88) = 72$$

$$P_2 = (20 - 0) = 20$$

$$P_3 = (28 - 0) + (88 - 48) + (125 - 108) = 85$$

$$P_4 = (48 - 0) + (108 - 68) = 88$$


- Average waiting time = $(72 + 20 + 85 + 88)/4 = 66\frac{1}{4}$
- Average completion time = $(125 + 28 + 153 + 112)/4 = 104\frac{1}{2}$
- Thus, Round-Robin Pros and Cons:
 - ▶ Better for short jobs, Fair (+)
 - ▶ Context-switching time adds up for long jobs (-)

Example of RR with Time Slice = 20s

- Example:

Thread	Burst Time (s)
P_1	53
P_2	8
P_3	68
P_4	24

- The Gantt chart is:



- Waiting time for

0 20 28 48 68 88 108 112 125 145 153

$$P_1 = (68 - 20) + (112 - 88) = 72$$

$$P_2 = (20 - 0) = 20$$

$$P_3 = (28 - 0) + (88 - 48) + (125 - 108) = 85$$

$$P_4 = (48 - 0) + (108 - 68) = 88$$


- Average waiting time = $(72 + 20 + 85 + 88)/4 = 66\frac{1}{4}$
- Average completion time = $(125 + 28 + 153 + 112)/4 = 104\frac{1}{2}$
- Thus, Round-Robin Pros and Cons:
 - ▶ Better for short jobs, Fair (+)
 - ▶ Context-switching time adds up for long jobs (-)

Example of RR with Time Slice = 20s

- Example:

Thread	Burst Time (s)
P_1	53
P_2	8
P_3	68
P_4	24

- The Gantt chart is:



- Waiting time for

$$P_1 = (68 - 20) + (112 - 88) = 72$$

$$P_2 = (20 - 0) = 20$$

$$P_3 = (28 - 0) + (88 - 48) + (125 - 108) = 85$$

$$P_4 = (48 - 0) + (108 - 68) = 88$$


- Average waiting time = $(72 + 20 + 85 + 88)/4 = 66\frac{1}{4}$
- Average completion time = $(125 + 28 + 153 + 112)/4 = 104\frac{1}{2}$
- Thus, Round-Robin Pros and Cons:
 - ▶ Better for short jobs, Fair (+)
 - ▶ Context-switching time adds up for long jobs (-)

Example of RR with Time Slice = 20s

- Example:

Thread	Burst Time (s)
P_1	53
P_2	8
P_3	68
P_4	24

- The Gantt chart is:



- Waiting time for

$$P_1 = (68 - 20) + (112 - 88) = 72$$

$$P_2 = (20 - 0) = 20$$

$$P_3 = (28 - 0) + (88 - 48) + (125 - 108) = 85$$

$$P_4 = (48 - 0) + (108 - 68) = 88$$


- Average waiting time = $(72 + 20 + 85 + 88)/4 = 66\frac{1}{4}$
- Average completion time = $(125 + 28 + 153 + 112)/4 = 104\frac{1}{2}$
- Thus, Round-Robin Pros and Cons:
 - ▶ Better for short jobs, Fair (+)
 - ▶ Context-switching time adds up for long jobs (-)

Example of RR with Time Slice = 20s

- Example:

Thread	Burst Time (s)
P_1	53
P_2	8
P_3	68
P_4	24

- The Gantt chart is:



- Waiting time for

0 20 28 48 68 88 108 112 125 145 153

$$P_1 = (68 - 20) + (112 - 88) = 72$$

$$P_2 = (20 - 0) = 20$$

$$P_3 = (28 - 0) + (88 - 48) + (125 - 108) = 85$$

$$P_4 = (48 - 0) + (108 - 68) = 88$$


- Average waiting time = $(72 + 20 + 85 + 88)/4 = 66\frac{1}{4}$
- Average completion time = $(125 + 28 + 153 + 112)/4 = 104\frac{1}{2}$
- Thus, Round-Robin Pros and Cons:
 - ▶ Better for short jobs, Fair (+)
 - ▶ Context-switching time adds up for long jobs (-)

Example of RR with Time Slice = 20s

- Example:

Thread	Burst Time (s)
P_1	53
P_2	8
P_3	68
P_4	24

- The Gantt chart is:



- Waiting time for

0 20 28 48 68 88 108 112 125 145 153

$$P_1 = (68 - 20) + (112 - 88) = 72$$

$$P_2 = (20 - 0) = 20$$

$$P_3 = (28 - 0) + (88 - 48) + (125 - 108) = 85$$

$$P_4 = (48 - 0) + (108 - 68) = 88$$

- Average waiting time = $(72 + 20 + 85 + 88)/4 = 66\frac{1}{4}$


- Average completion time = $(125 + 28 + 153 + 112)/4 = 104\frac{1}{2}$

- Thus, Round-Robin Pros and Cons:

- ▶ Better for short jobs, Fair (+)
- ▶ Context-switching time adds up for long jobs (-)

Round-Robin Discussion

- How do you choose time slice?
 - ▶ What if too big?
 - ★ Response time suffers
 - ▶ What if infinite (∞)?
 - ★ Get back FCFS
 - ▶ What if time slice too small?
 - ★ Throughput suffers!
- Actual choices of time slice:
 - ▶ In practice, need to balance short-job performance and long-job throughput:
 - ★ Typical time slice today is between **10ms – 100ms**
 - ★ Typical context-switching overhead is **0.1ms – 1ms**
 - ★ Roughly **1%** overhead due to context-switching



Comparisons between FCFS and RR

- Assuming zero-cost context-switching time, is RR always better than FCFS?
- Simple example:
10 jobs, each takes 100s of CPU time RR time slice of 1s.
All jobs start at the same time.
- Completion times:

Job #	FIFO	RR
1	100	991
2	200	992
...
9	900	999
10	1000	1000

- Both RR and FCFS finish at the same time
 - Average response time is much worse under RR!
- Cache state must be shared between all jobs with RR but can be devoted to each job with FIFO

FCFS vs. Round Robin


- FCFS Scheduling:

- ▶ Run threads to completion in order of submission
- ▶ Pros: Simple
- ▶ Cons: Short jobs get stuck behind long ones

- Round-Robin Scheduling:

- ▶ Give each thread a small amount of CPU time when it executes; cycle between all ready threads
- ▶ Pros: Better for short jobs
- ▶ Cons: Poor when jobs have same length

Earlier Example with Different Time Slices



	Time slice	P_1	P_2	P_3	P_4	Average
Wait Time	Best FCFS	32	0	85	8	$31\frac{1}{4}$
	Q = 1	84	22	85	57	62
	Q = 5	82	20	85	58	$61\frac{1}{4}$
	Q = 8	80	8	85	56	$57\frac{1}{4}$
	Q = 10	82	10	85	68	$61\frac{1}{4}$
	Q = 20	72	20	85	88	$66\frac{1}{4}$
	Worst FCFS	68	145	0	121	$83\frac{1}{2}$
Completion Time	Best FCFS	85	8	153	32	$69\frac{1}{2}$
	Q = 1	137	30	153	81	$100\frac{1}{2}$
	Q = 5	135	28	153	82	$99\frac{1}{2}$
	Q = 8	133	16	153	80	$95\frac{1}{2}$
	Q = 10	135	18	153	92	$99\frac{1}{2}$
	Q = 20	125	28	153	112	$104\frac{1}{2}$
	Worst FCFS	121	153	68	145	$121\frac{3}{4}$

SJF and SRTF

Assumption: job lengths are known in advance.


- Shortest Job First (SJF)
 - ▶ Non-preemptive
- Shortest Remaining Time First (SRTF)
 - ▶ Preemptive version
- Example
 - ▶ $P1 = 6\text{sec}$, $P2 = 8\text{sec}$, $P3 = 7\text{sec}$, $P4 = 3\text{sec}$
 - ▶ All arrive at the same time



- Can you do better than SRTF in terms of average response time?

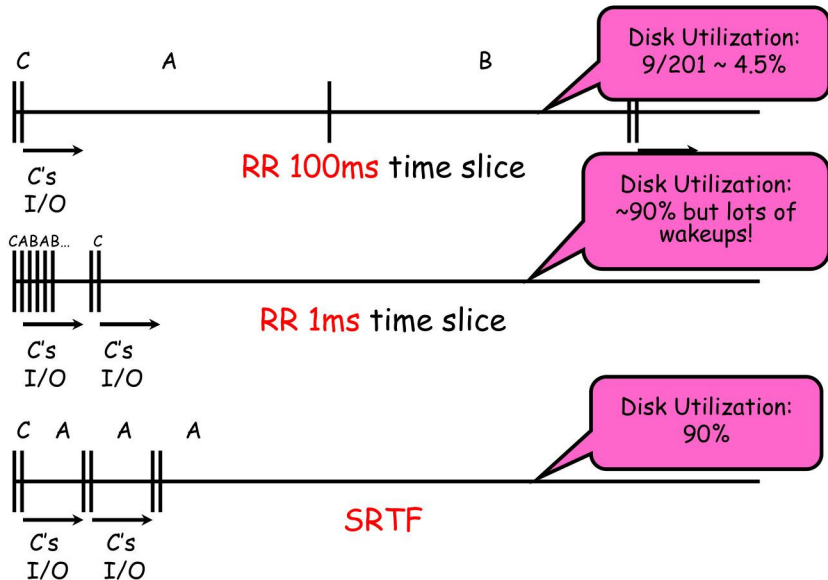
- SJF/SRTF \Rightarrow best at minimizing average response time
 - ▶ Provably optimal:
 - ★ SJF among non-preemptive
 - ★ SRTF among preemptive
 - ▶ SRTF at least as good as SJF \Rightarrow focus on SRTF
- Comparison of SRTF with FCFS and RR
 - ▶ What if all jobs the same length?
 - ★ SRTF becomes the same as FCFS
i.e. FCFS is best can do if all jobs the same length
 - ▶ What if jobs have varying length?
 - ★ SRTF vs RR \Rightarrow short jobs not stuck behind long ones

Example to illustrate benefits of SRTF




- Three jobs:
 - ▶ A,B: both CPU bound, run for 1 week
⇒ A or B alone use 100% of the CPU
 - ▶ C: I/O bound, loop 1ms CPU, 9ms disk I/O
⇒ C alone uses 90% of the disk
- With FCFS/FIFO:
 - ▶ Once A or B get in, keep CPU for two weeks
- What about RR or SRTF?

SRTF Example continued:



Multi-Level Feedback Scheduling



- A method for exploiting past behavior
 - ▶ **Multiple queues, each with different priority**
 - ★ Higher priority queues often considered "foreground" tasks
 - ▶ **Each queue has its own scheduling algorithm**
 - ★ e.g. foreground – RR, background – FCFS
 - ★ Sometimes multiple RR priorities with quantum increasing exponentially (highest:1ms, next:2ms, next: 4ms, etc)
- Adjust each job's priority as follows (details vary)
 - ▶ Job starts in highest priority queue
 - ▶ If timeout expires, drop one level
 - ▶ If timeout doesn't expire, push up one level (or to top)

Scheduling Details

- Result approximates SRTF:
 - ▶ CPU bound jobs drop like a rock
 - ▶ Short-running I/O bound jobs stay near top
- Scheduling must be done between the queues
 - ▶ **Fixed priority scheduling:**
 - ★ Serve all from highest priority, then next priority, etc.
 - ▶ **Time slice:**
 - ★ Each queue gets a certain amount of CPU time
 - ★ E.g., 70% to highest, 20% next, 10% lowest
- **Countermeasure:** user action that can foil intent of the OS designer
 - ▶ For multilevel feedback, put in a bunch of meaningless I/O to keep job's priority high
 - ▶ Of course, if everyone did this, wouldn't work!

Priority Scheduling

- Obvious
 - ▶ Not all threads are equal, so rank them
- The method
 - ▶ Assign each thread a priority
 - ▶ Run the thread with highest priority in the ready queue first
 - ▶ Adjust priority dynamically
- Why adjust priorities dynamically?
 - ▶ Problem: **Starvation**
 - ▶ Solution: **Aging**

Lottery Scheduling

- Motivations
 - ▶ SRTF does well with average response time, but unfair
- Lottery method
 - ▶ Give each job a number of tickets
 - ▶ Randomly pick a winning ticket
 - ▶ To approximate SRTF, give short jobs more tickets
 - ▶ To avoid starvation, give each job at least one ticket
 - ▶ Cooperative processes can exchange tickets
- Question: how do you compare this method with priority scheduling?
 - ▶ Behaves gracefully as load changes

Lottery Scheduling Example

- Assume short jobs get 10 tickets, long jobs get 1 ticket

# short jobs/ # long jobs	% of CPU each short jobs get	% of CPU each long jobs get
1/1	91%	9%
0/2	N/A	50%
2/0	50%	N/A
10/1	9.9%	0.99%
1/10	50%	5%

Summary

- Different scheduling goals
 - ▶ Depend on what systems you build
- Scheduling algorithms
 - ▶ FCFS is simple and good for jobs with the same length
 - ▶ RR is good for short jobs with different length
 - ▶ SJF and SRTF (theoretically) give the optimal average response time
 - ▶ Priority scheduling and its variations (e.g., multi-level feedback scheduling) are most common in practice
 - ▶ Lottery scheduling is flexible

References

- A. S. Tanenbaum, Modern Operating Systems.
- A. Silberschatz et. al., Operating System Concepts.

Thanks for your attention!

Questions?