

Message Passing

Loïc Guégan

Adapted from P. Ha @ UiT, K. Li @ Princeton,
A. S. Tanenbaum @ 2008, A. Silberschatz @ 2009 and B. Wilkinson @ 2004

UiT The Arctic University of Norway


Spring - 2024

- What is message passing ?
 - ▶ Semantics
 - ▶ How to use
- How to implement message passing?
 - ▶ Implementation issues
- Examples of message passing systems

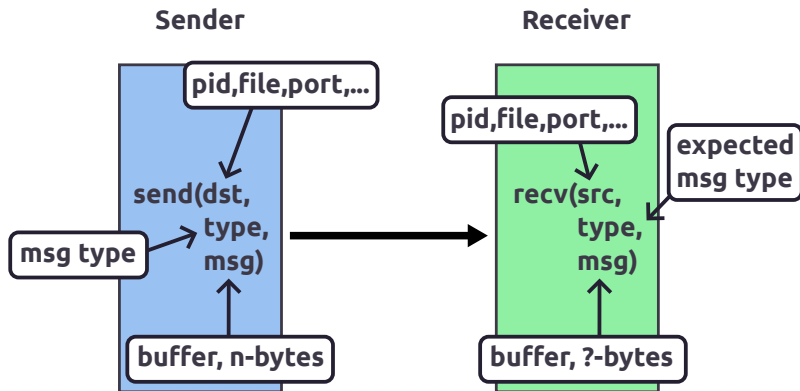
Overview of communication mechanisms

Programs	Concurrent Programs			
Higher-level API	Shared Variables Locks Semaphores Monitors		Message Passing Send/Receive	
Hardware	Load/Store	Disable Ints	Test&Set	Comp&Swap

Big picture




Send and Receive primitives




Many ways to design the message passing API

Synchronous Message Passing

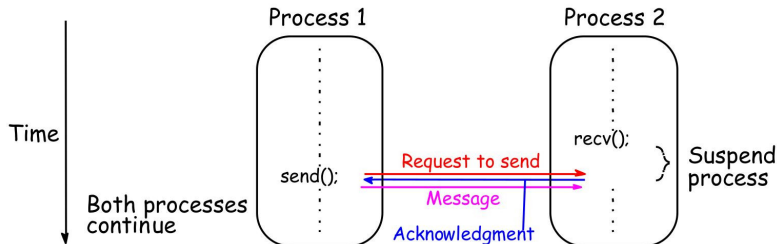
- Synchronizing processes
 - ▶ Sender: signal the receiver process that a particular event happens
 - ▶ Receiver: blocked until the event has happened
- Moving data between processes
 - ▶ Sender: send data to the receiver when data is ready
 - ▶ Receiver: collect data when the data has arrived and process is ready



Synchronous send() and recv() using 3-way protocol



(a) When `send()` occurs before `recv()`



(b) When `recv()` occurs before `send()`

Example: Producer-Consumer

```
Producer() {
```

```
    ...
```

```
    while (1) {
```

```
        produce item;
```

```
        recv(Consumer, &credit);
```

```
        send(Consumer, item);
```

```
    }
```

```
}
```

```
Consumer() {
```

```
    ...
```

```
    for (i=0; i<N; i++)
```

```
        send(Producer, credit);
```

```
        while (1) {
```

```
            recv(Producer, &item);
```

```
            send(Producer, credit);
```

```
            consume item;
```

```
        }
```

```
}
```

• Questions

- ▶ Would it work with multiple producers and 1 consumer?
- ▶ Would it work with 1 producer and multiple consumers?
- ▶ What about multiple producers and multiple consumers?

Summary


- Mechanism for processes to communicate and to synchronize their actions
- Allows to avoid using shared variables
- If P and Q wish to communicate:
 - ▶ Establish a communication link between them
 - ▶ Exchange messages via **send/receive**
- Layers:
 - ① physical (e.g. main memory, hardware bus)
 - ② **logical (e.g. logical properties)**
 - ★ Direct vs. indirect communication
 - ★ Synchronous vs. asynchronous communication
 - ★ Automatic vs. explicit buffering

Outline

- Message Passing
- Implementation issues
- Examples of message-passing systems

Implementation issues

- Buffering messages
- Direct vs. indirect
- Asynchronous vs. synchronous
- How to handle exceptions?



Buffering messages

- No buffering

- ▶ Sender must wait until the receiver receives the message
- ▶ Rendezvous on each message


- Bounded buffer

- ▶ Finite size
- ▶ Sender is blocked on buffer full



- Unbounded buffer

- ▶ "Infinite" size
- ▶ Sender is never blocked




Direct communication

- Processes must name each other explicitly:
 - ▶ **send**(P,message) - send a message to process P
 - ▶ **receive**(Q,message) - receive a message from process Q
- Properties of communication link
 - ▶ A link is associated with exactly one pair of communicating processes
 - ▶ Links are established automatically


Direct communication with buffer

- A single buffer at the receiver
 - ▶ More than one process may send messages to the receiver
 - ▶ `recv(ANY,msg)`
 - ▶ To broadcast a message, a sender must send separate messages to receivers
- A buffer at each sender
 - ▶ A sender may send a message to multiple receivers
 - ▶ To get a message, it requires searching through the whole buffer
- What if process id/name changes?



Indirect communication

- Use mailbox as the abstraction
 - ▶ **send**(A, message) - send a message to mailbox A
 - ▶ **recv**(A, message) - receive a message from mailbox A
 - ▶ Require create/destroy a mailbox
 - ▶ Allow many-to-many communication
- Mailbox sharing
 - ▶ P1, P2, and P3 share mailbox A. P1 sends, P2 and P3 receive.
 - ▶ Who gets the message?
- Mailbox vs. pipe
 - ▶ A mailbox allows many to many communication
 - ▶ A pipe implies one sender and one receiver




- Message Passing
- Implementation issues
 - ▶ Buffering
 - ▶ Indirection
 - ▶ Blocking vs. nonblocking
 - ▶ Exceptions

Blocking vs. Nonblocking: Send

• Blocking

- ▶ Initiate data transfer
- ▶ Either:
 - 1 Block until data is out of its source memory (buffering)
 - 2 Block until data is received (no buffering)

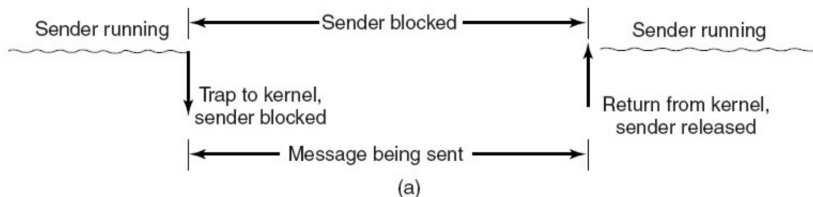


• Nonblocking

- ▶ Initiate data transfer and return
- ▶ Completion
 - ★ Require applications to check status
 - ★ Notify or signal the application

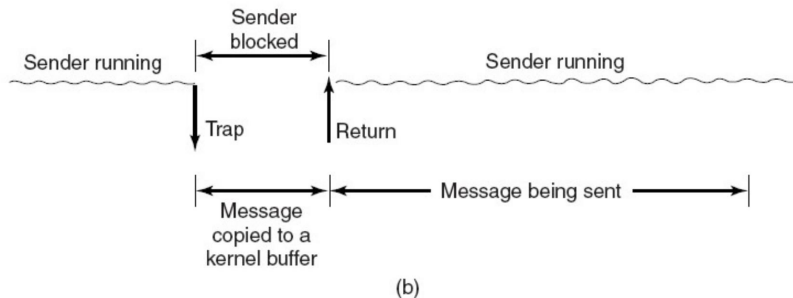
```
status = async_send( dest, type, msg )  
...  
if !send_complete( status )  
    wait for completion;  
...  
use msg data structure;  
...
```

Blocking versus Nonblocking: Send (2)



A blocking send call

Blocking versus Nonblocking: Send (3)



A nonblocking send call

Blocking vs. Nonblocking: Receive

- Blocking
 - ▶ Block until a message is available

- Nonblocking
 - ▶ Return either a valid message or a null
 - ▶ Null: probe(src)

msg transfer resource

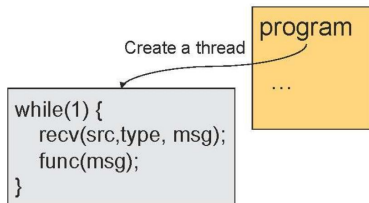
→ `recv(src, type, msg)`

```
status = async_recv( src, type, msg );  
if ( status == SUCCESS )  
    consume msg;  
...  
while ( probe(src) != HaveMSG )  
    wait for msg arrival  
recv( src, type, msg );  
consume msg;
```

Event Handler vs. Receive


- `hrecv(src, type, msg, func)`
 - ▶ `msg` is an arg of `func`
 - ▶ Execute "`func`" on a message arrival
- Which one is more powerful?
 - ▶ Recv with a thread can emulate a Handler
 - ▶ Handler can be used to emulate recv by using Monitor

```
void func( char * msg ) {  
    ...  
}  
  
...  
hrecv( src, type, msg, func)  
...
```



Example: Keyboard Input

- How do you implement keyboard input?
 - ▶ Need an interrupt handler
 - ▶ Generate a mbox message from the interrupt handler
- Suppose a keyboard device thread converts input characters into an mbox message
 - ▶ How would you synchronize between the keyboard interrupt handler and device thread?
 - ▶ How can a device thread convert input into mbox messages?



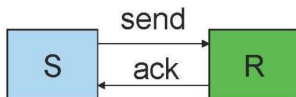
- Message Passing
- Implementation issues
 - ▶ Buffering
 - ▶ Indirection
 - ▶ Blocking vs. nonblocking
 - ▶ Exceptions

Exception: Process Termination

- R waits for a message from S, but S has terminated
 - ▶ Problem: R may be blocked forever
- S sends a message to R, but R has terminated
 - ▶ Problem: S has no buffer and will be blocked forever




Exception: Message Loss




- Use ack and timeout to detect and retransmit a lost message
 - ▶ Receiver sends an ack message for each message
 - ▶ Sender is blocked until an ack message is back or timeout
`status = send(dest, msg, timeout);`
 - ▶ If timeout happens and no ack, sender retransmits the message
- Issues
 - ▶ Losing ack messages
 - ▶ Duplicates

Exception: Message Loss (2)

- Retransmission must handle
 - ▶ Duplicate messages on receiver side
 - ▶ Out-of-sequence ack messages on sender side
- Retransmission
 - ▶ Use sequence number for each message to identify duplicates
 - ▶ Remove duplicates on receiver side
 - ▶ Sender retransmits on an out-of-sequence ack
- Reduce ack messages
 - ▶ Bundle ack messages
 - ▶ Receiver sends NOACK messages: can be complex
 - ▶ Piggy-back acks in send messages



Exception: Message Corruption



- Detection

- ▶ Compute a checksum over the entire message and send the checksum (e.g. CRC) as part of the message
- ▶ Recompute a checksum at receiver and compare with the checksum in the message

- Correction

- ▶ Trigger retransmission
- ▶ Use correction codes (e.g. ECC) to recover

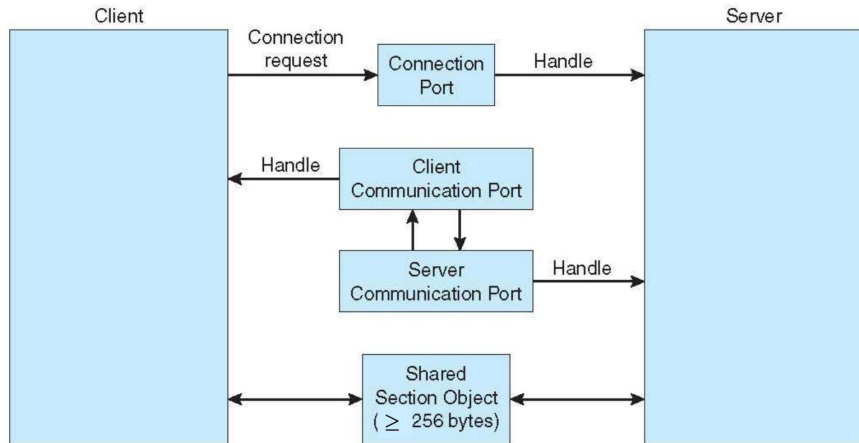
Outline

- Message Passing
- Implementation issues
- Examples of message passing systems

Mach microkernel (Mac OS X)

- Mach communication is message based
 - ▶ System calls are messages
 - ▶ Each task gets two mailboxes at creation- Kernel and Notify
 - ▶ Only three system calls needed for message transfer
`msg_send()`, `msg_receive()`, `msg_rpc()`
 - ▶ Mailboxes needed for communication, created via
`port_allocate()`

Local Procedure Calls in Windows XP



References

- A. S. Tanenbaum, Modern Operating Systems.
- A. Silberschatz et. al., Operating System Concepts.
- B. Wilkinson et. al., Parallel Programming: Techniques and Applications Using Networked Workstations and Parallel Computers.
- B. Barney, "Message Passing Interface (MPI)", Livermore Computing.

Thanks for your attention!

Questions?