

Virtual Memory, Address-Translation and Paging

INF-2201 Operating Systems – Spring 2024

Loïc Guégan (loic.guegan@uit.no)

Based on presentations created by:

**Issam Raïs, Lars Ailo Bongo, Tore Brox-Larsen, Bård Fjukstad,
Daniel Stødle, Otto Anshus, Thomas Plagemann, Kai Li, Andy
Bavier**

Tanenbaum & Bo, Modern Operating Systems: 4th ed.

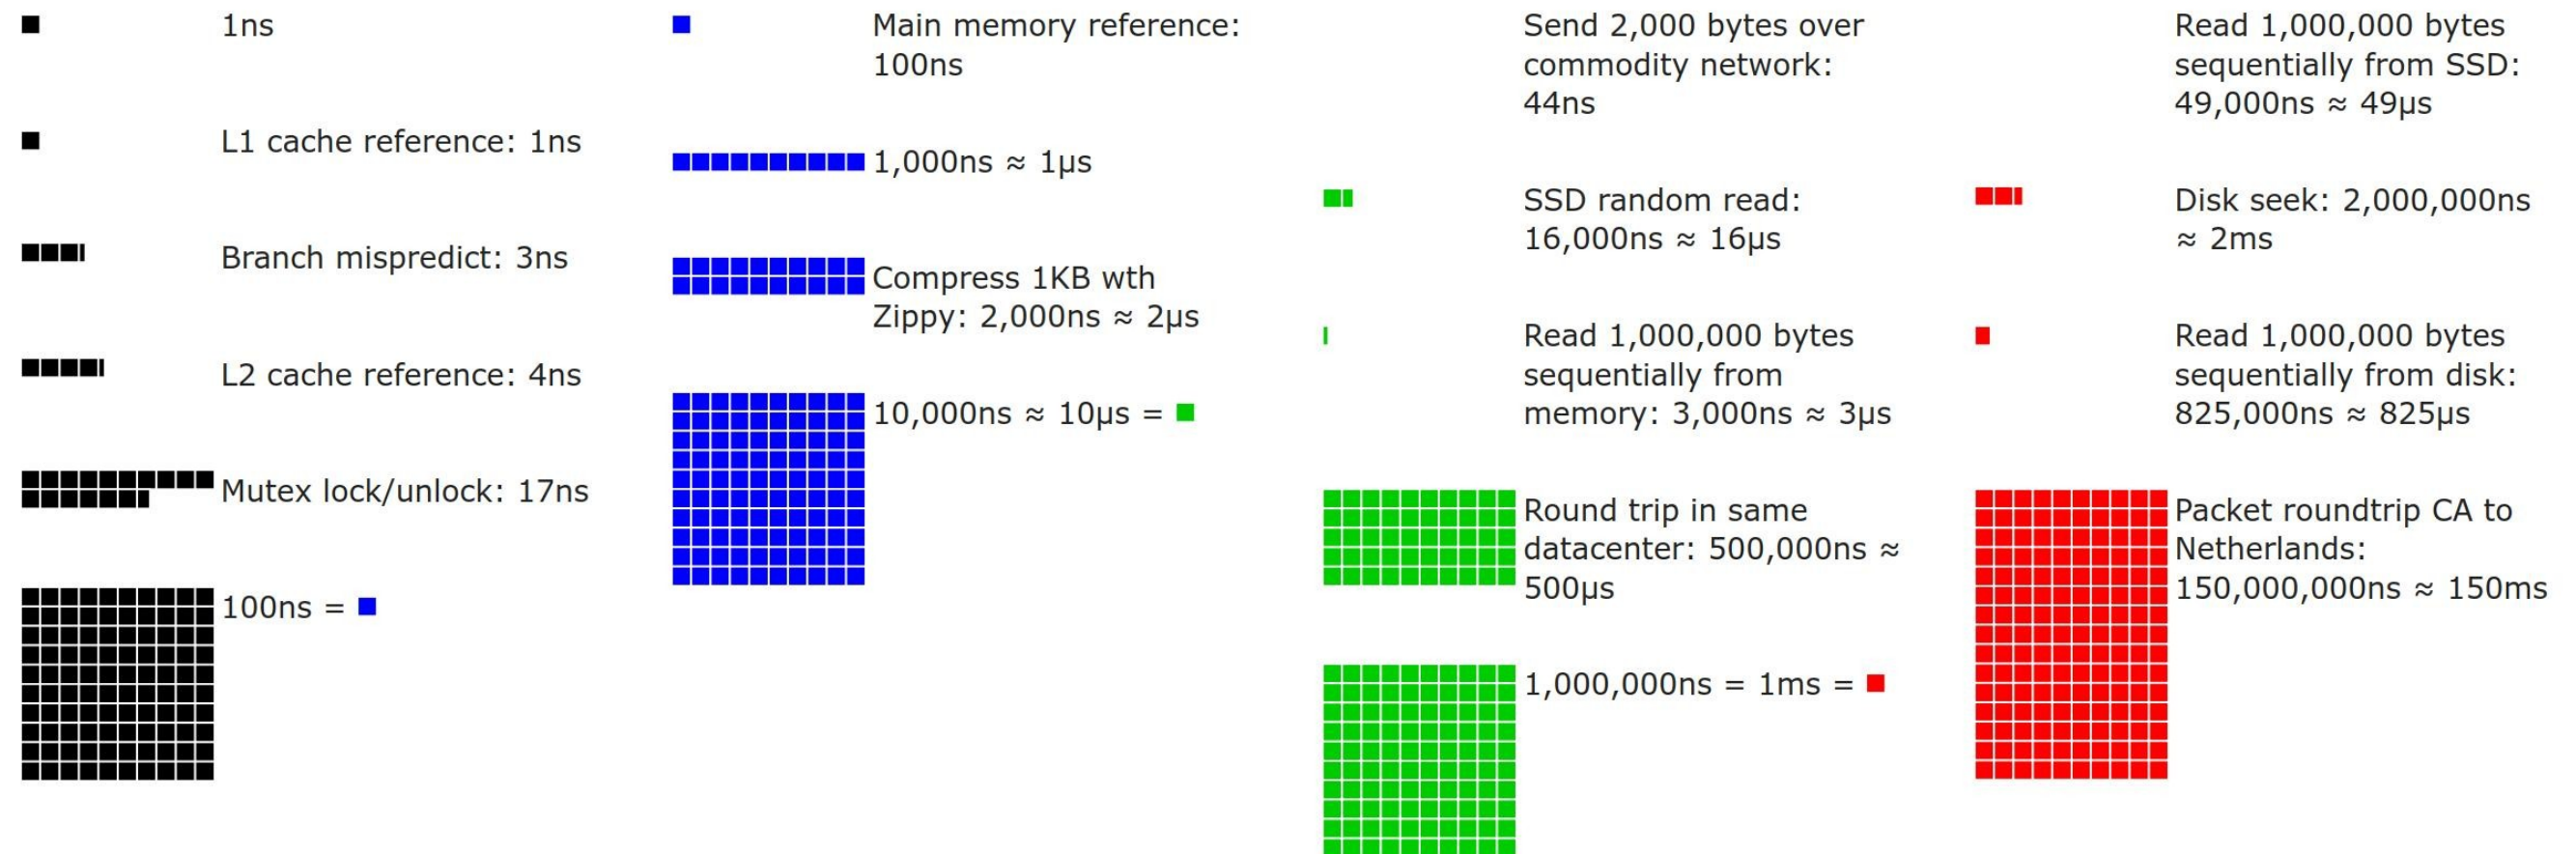
Overview

- Part 1
 - Virtual Memory
 - Virtualization
 - Protection
 - Address Translation
 - Base and bound
 - Segmentation
 - Paging
 - Translation look-ahead buffer
- Part 2: Paging and replacement
- Part 3: Design Issues

Latency Numbers Every Programmer Should Know

Latency Numbers Every Programmer Should Know


2020



https://people.eecs.berkeley.edu/~rcs/research/interactive_latency.html

The Big Picture

- DRAM is fast, but relatively expensive
 - \$25/GB
 - 20-30ns latency
 - 10-80GB's/sec
- Disk is inexpensive, but slow
 - \$0.2-1/GB (100 less expensive)
 - 5-10ms latency (100K times slower)
 - 40-80MB/sec per disk (1,000 times less)
- Our goals
 - Run programs as efficiently as possible
 - Make the system as safe as possible



Issues

- Many processes
 - The more processes a system can handle, the better
- Address space size
 - Many small processes whose total size may exceed memory
 - Even one process may exceed the physical memory size
- Protection
 - A user process should not crash the system
 - A user process should not do bad things to other processes

No Memory Abstraction

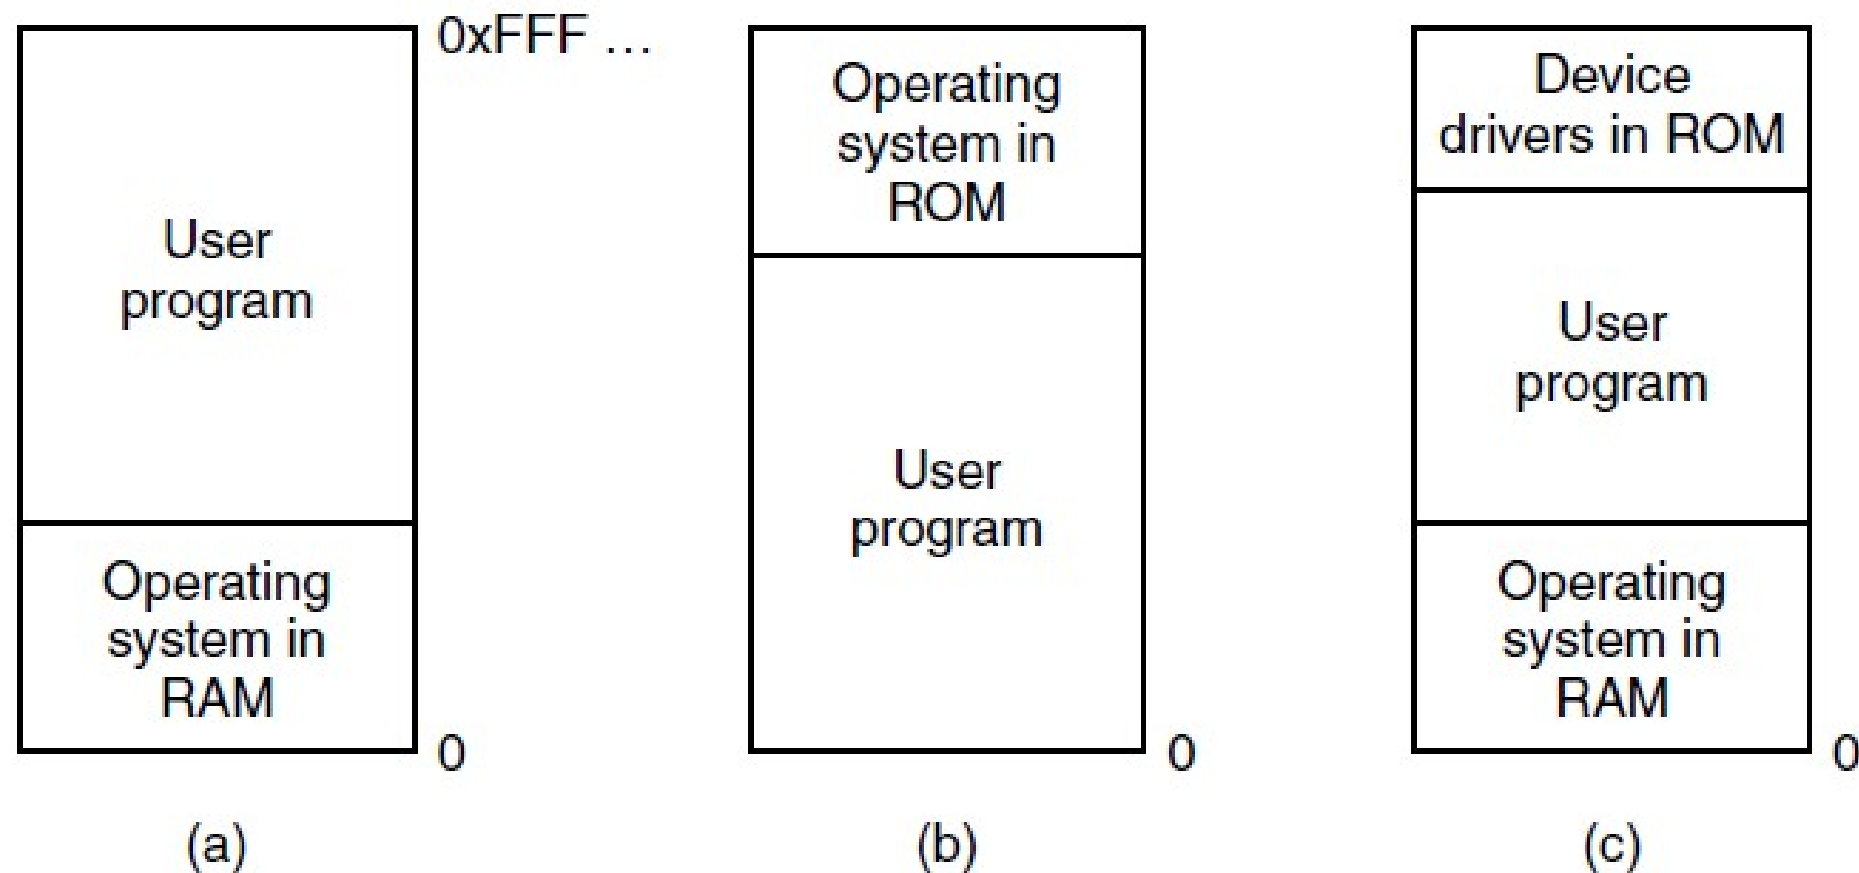
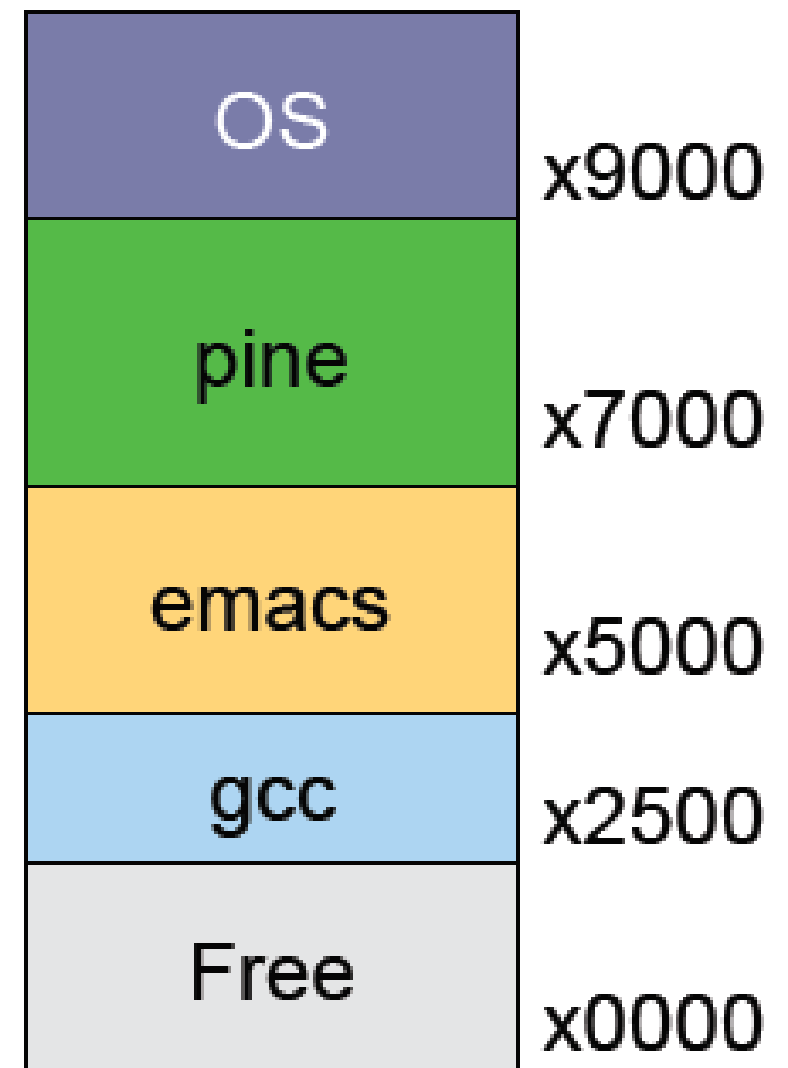


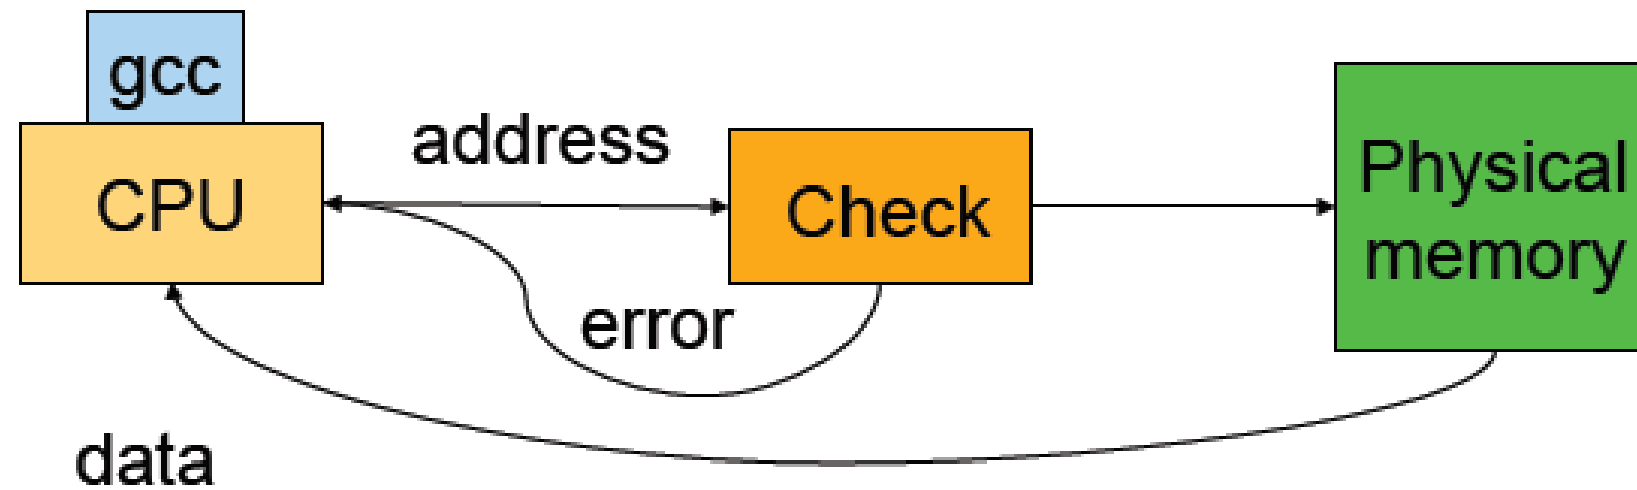
Figure 3-1. Three simple ways of organizing memory with an operating system and one user process. Other possibilities also exist

A Simple System

- Only physical memory
 - Applications use physical memory directly
- Run three processes
 - emacs, pine, gcc
- What if
 - gcc has an address error?
 - emacs writes at x7050?
 - pine needs to expand?
 - emacs needs more memory than is on the machine?




Protection Issues



- Errors in one process should not affect others
- For each process, check each load and store instruction to allow only legal memory references

Expansion or Transparency Issue

- A process should be able to run regardless of its physical location or the physical memory size
- Give each process a large, static “fake” address space
- As a process runs, relocate each load and store to its actual memory



Virtual Memory


- Flexible
 - Processes can move in memory as they execute, partially in memory and partially on disk
- Simple
 - Make applications very simple in terms of memory accesses
- Efficient
 - 20/80 rule: 20% of memory gets 80% of references
 - Keep the 20% in physical memory
- Design issues
 - How is protection enforced?
 - How are processes relocated?
 - How is memory partitioned?

Address Mapping and Granularity

- Must have some “mapping” mechanism
 - Virtual addresses map to DRAM physical addresses or disk addresses
- Mapping must have some granularity
 - Granularity determines flexibility
 - Finer granularity requires more mapping information
- Extremes
 - Any byte to any byte: mapping equals program size
 - Map whole segments: larger segments problematic

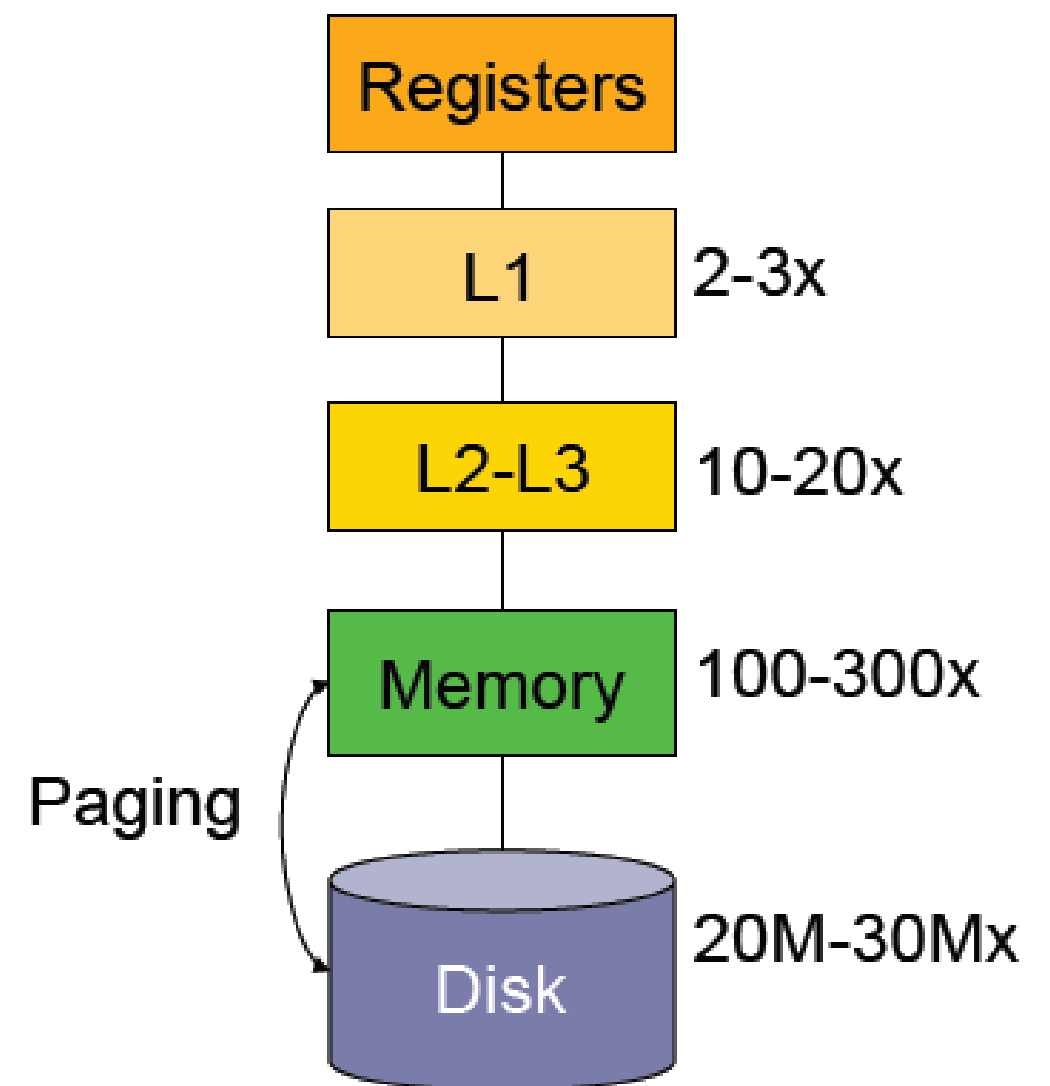
Generic Address Translation

- Memory Management Unit (MMU) translates virtual address into physical address for each load and store
- Software (privileged) controls the translation
- CPU view
 - Virtual addresses
- Each process has its own memory space [0, high]
 - Address space
- Memory or I/O device view
 - Physical addresses



Goals of Translation


- Implicit translation for each memory reference
- A hit should be very fast
- Trigger an exception on a miss
- Protected from user's faults



Base and Bound




- Built in Cray-1
- Each process has a pair (base, bound)
- Protection
 - A process can only access physical memory in [base, base+bound]
- On a context switch
 - Save/restore base, bound registers
- Pros
 - Relocation possible
 - Simple
 - Flat and no paging
- Cons
 - Fragmentation
 - Hard to share
 - Difficult to use disks




Segmentation

- Each process has a table of (seg, size)
- Treats (seg, size) as a fine-grained (base, bound)
- Protection
 - Each entry has (nil, read, write, exec)
- On a context switch
 - Save/restore the table and a pointer to the table in kernel memory
- Pros
 - Efficient
 - Easy to share
- Cons
 - Segments must be contiguous in physical memory
 - External Fragmentation



Paging

- Use a fixed size unit called page instead of segment
- Use a page table to translate
 - “Principle: Introducing (one level of) indirection”
- Various bits in each entry
- Context switch
 - Similar to segmentation
- What should page size be?
- Pros
 - Simple allocation
 - Easy to share
- Cons
 - Big table
 - Fragmentation in table
- How to deal with holes?



Paging (1)

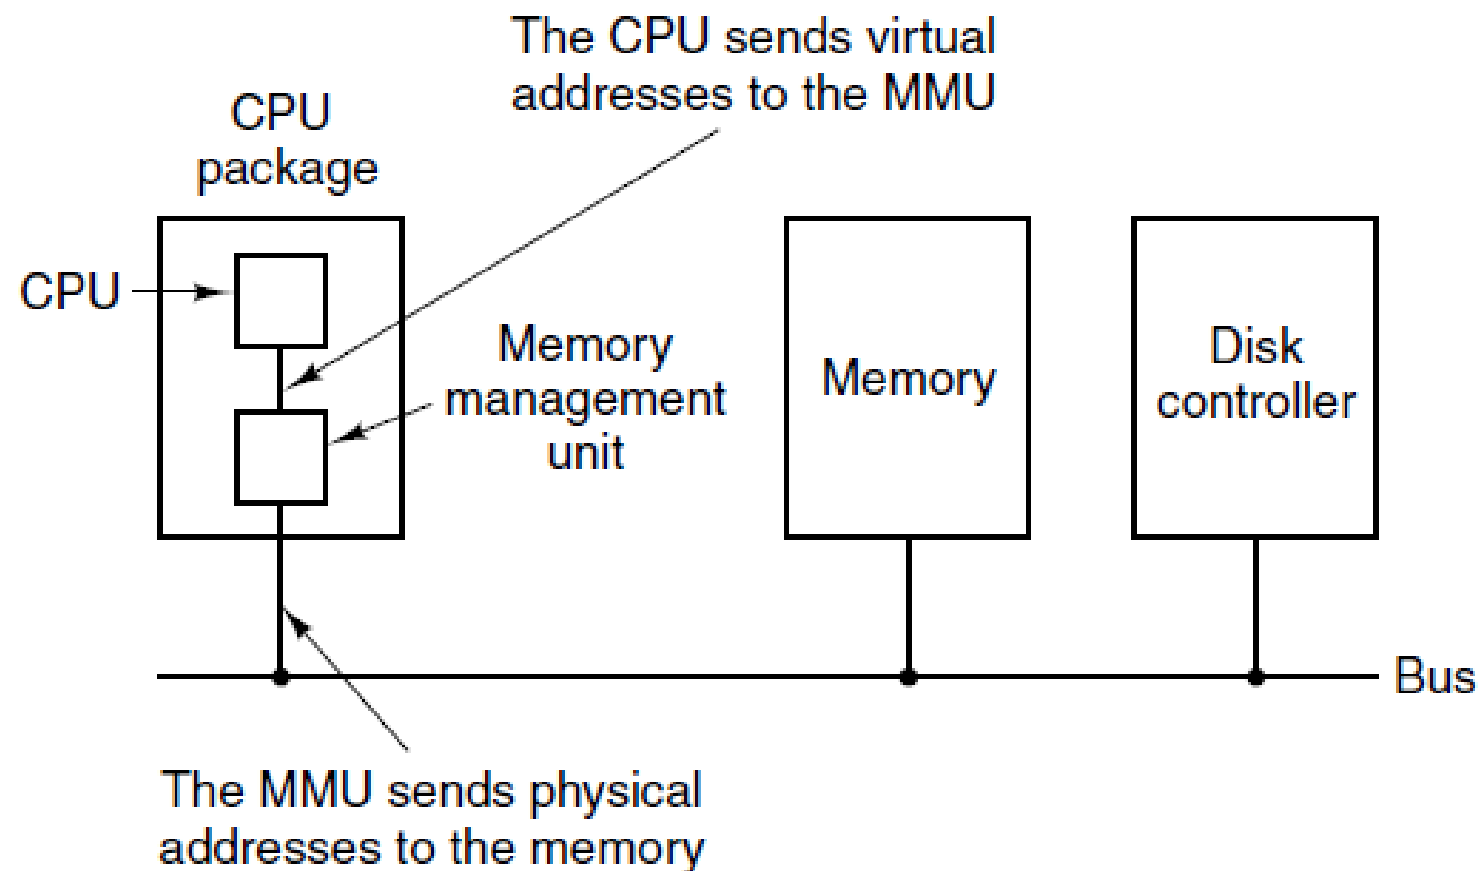


Figure 3-8. The position and function of the MMU. Here the MMU is shown as being a part of the CPU chip because it commonly is nowadays. However, logically it could be a separate chip and was years ago.

Paging (2)

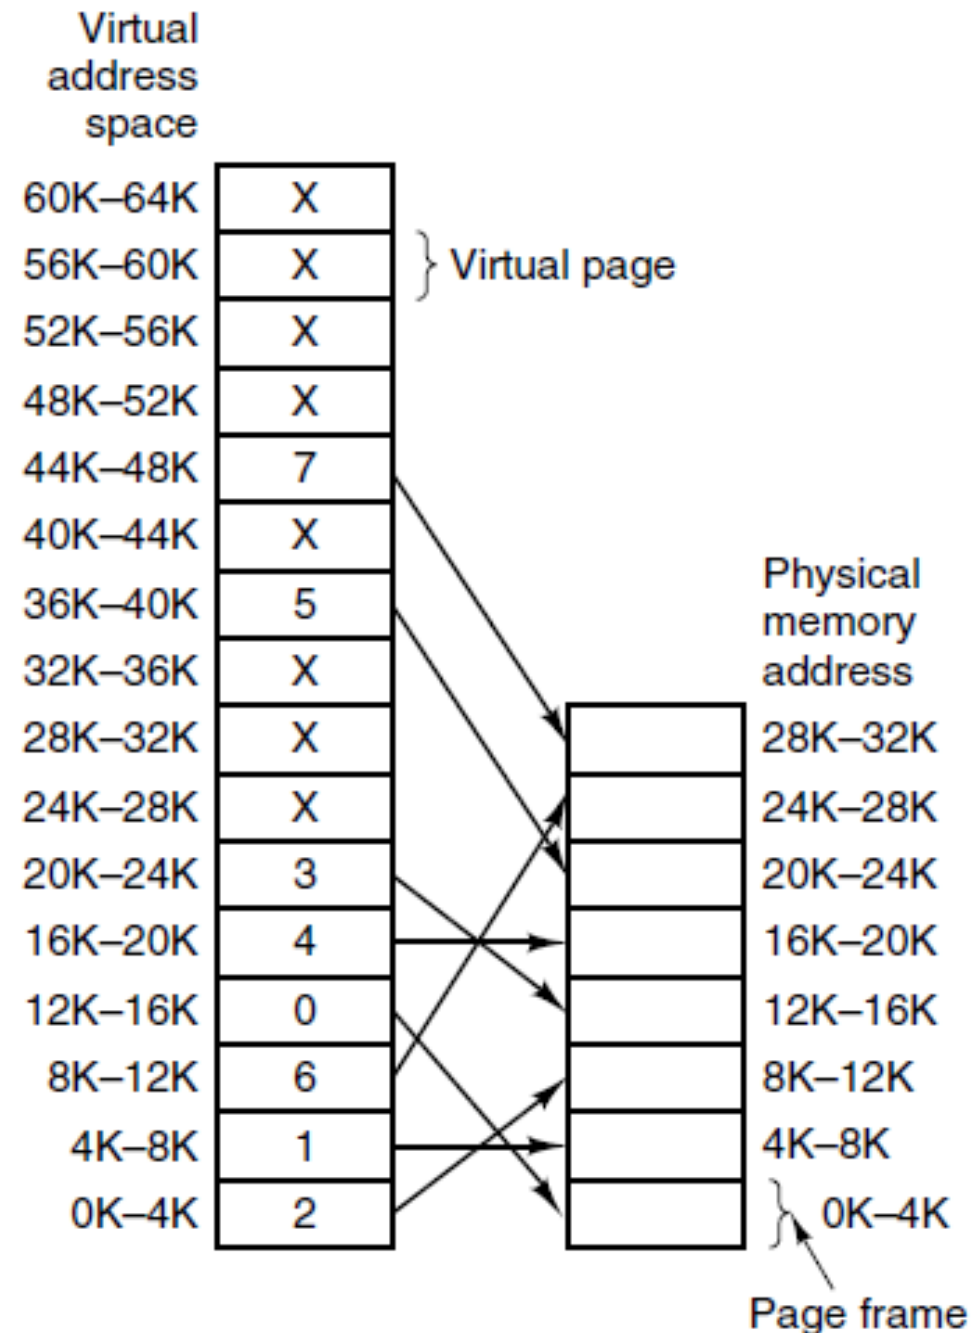



Figure 3-9. The relation between virtual addresses and physical memory addresses is given by the page table. Every page begins on a multiple of 4096 and ends 4095 addresses higher, so 4K-8K really means 4096-8191 and 8K to 12K means 8192-12287

Paging (3)

Figure 3-10. The internal operation of the MMU with 16 4-KB pages.



Structure of a Page Table Entry

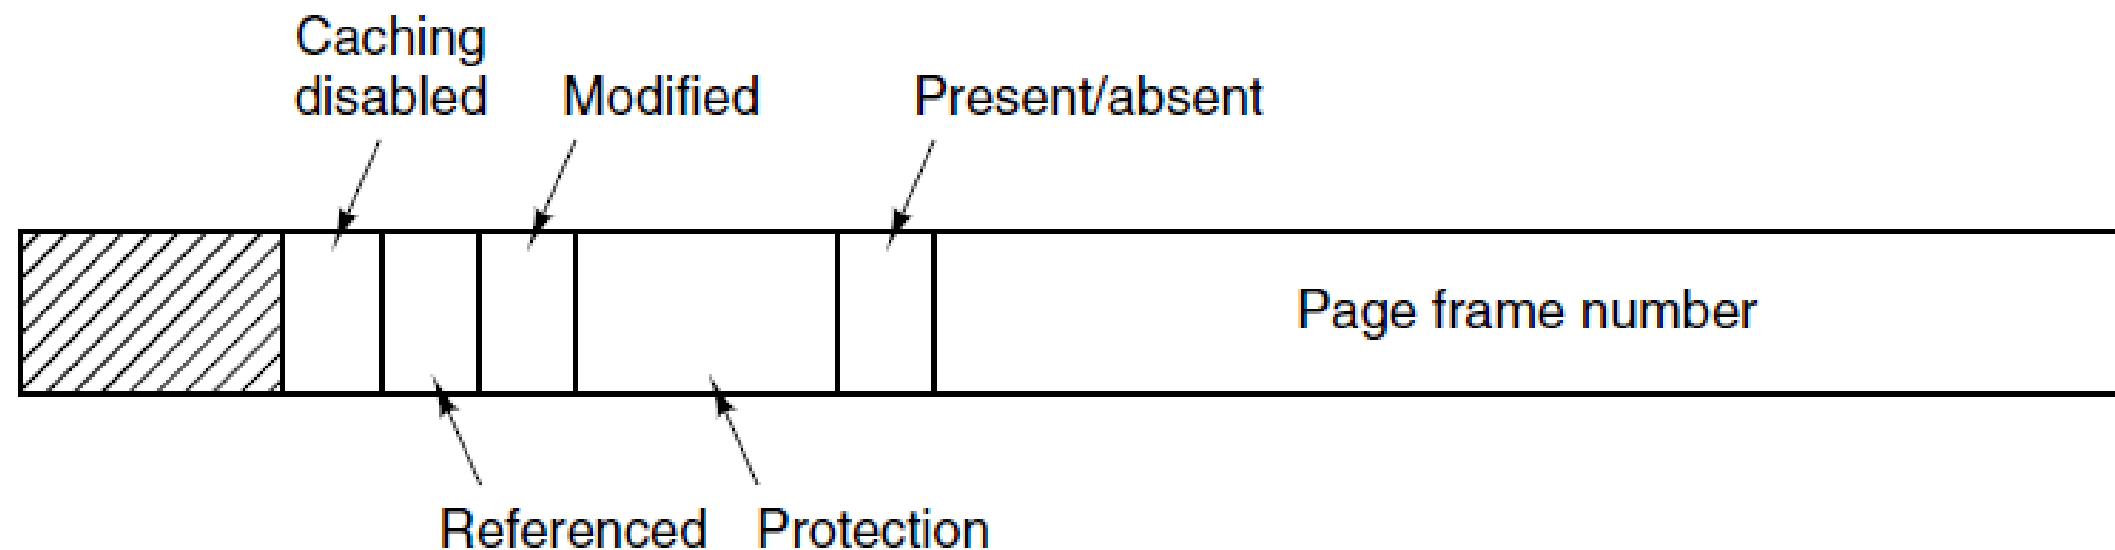


Figure 3-11. A typical page table entry.

Speeding Up Paging

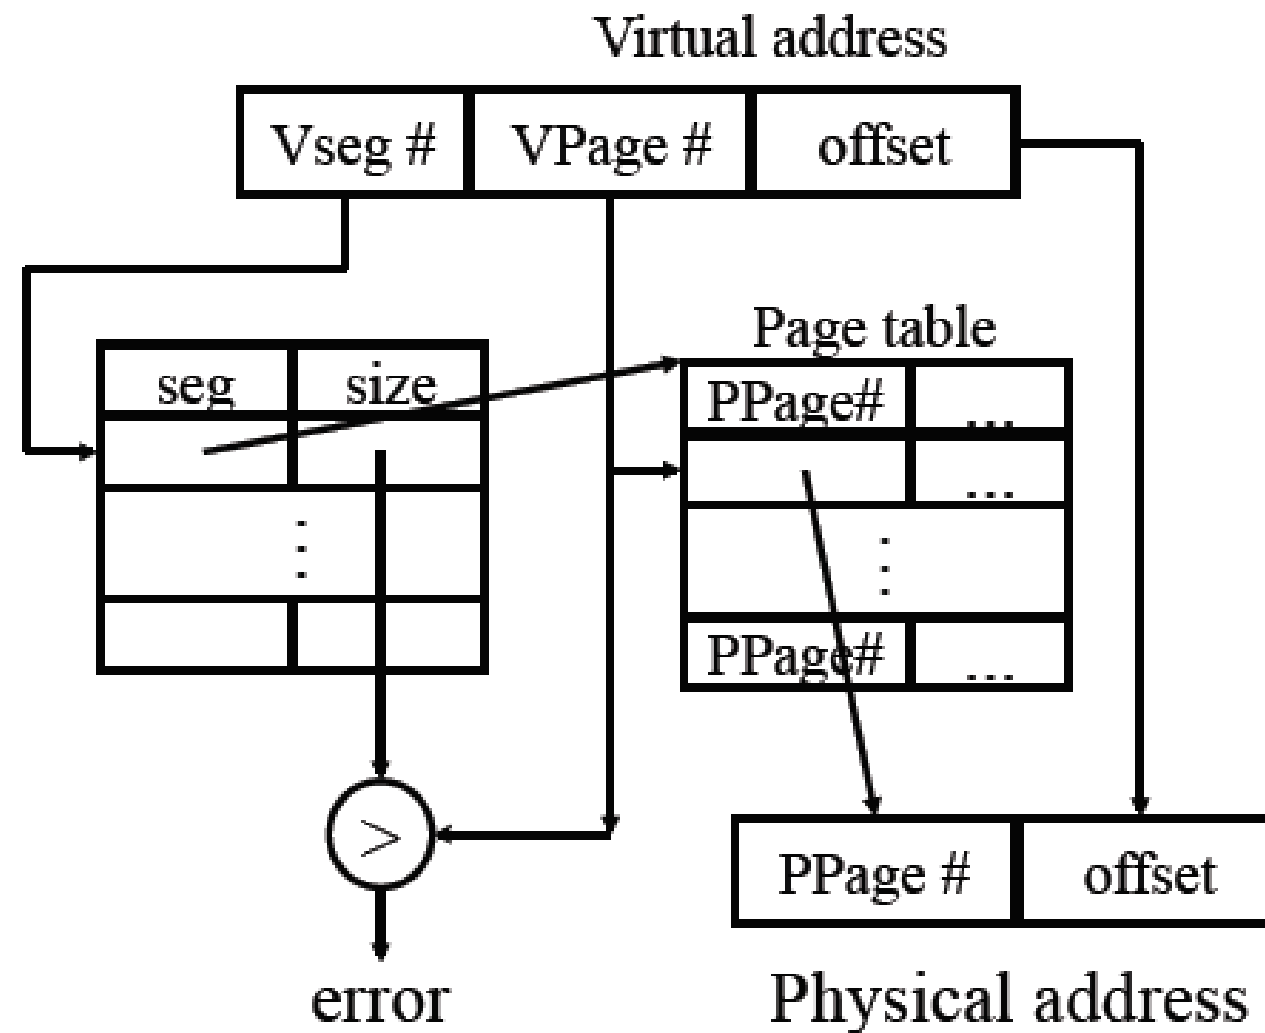
Major issues faced:

- 1.The mapping from virtual address to physical address must be fast.
- 2.If the virtual address space is large, the page table will be large.


How Many PTEs Do We Need?

- Assume 4KB page
 - Equals “low order” 12 bits
- Worst case for 32-bit address machine
 - # of processes $\times 2^{20}$
 - 2^{20} PTEs per page table (~ 4 Mbytes), but there might be 10K processes. They won't fit in memory together
- What about 64-bit address machine?
 - # of processes $\times 2^{52}$
 - A page table cannot fit in a disk (2^{52} PTEs = 16PBytes)!

Segmentation with Paging



Multiple-Level Page Tables

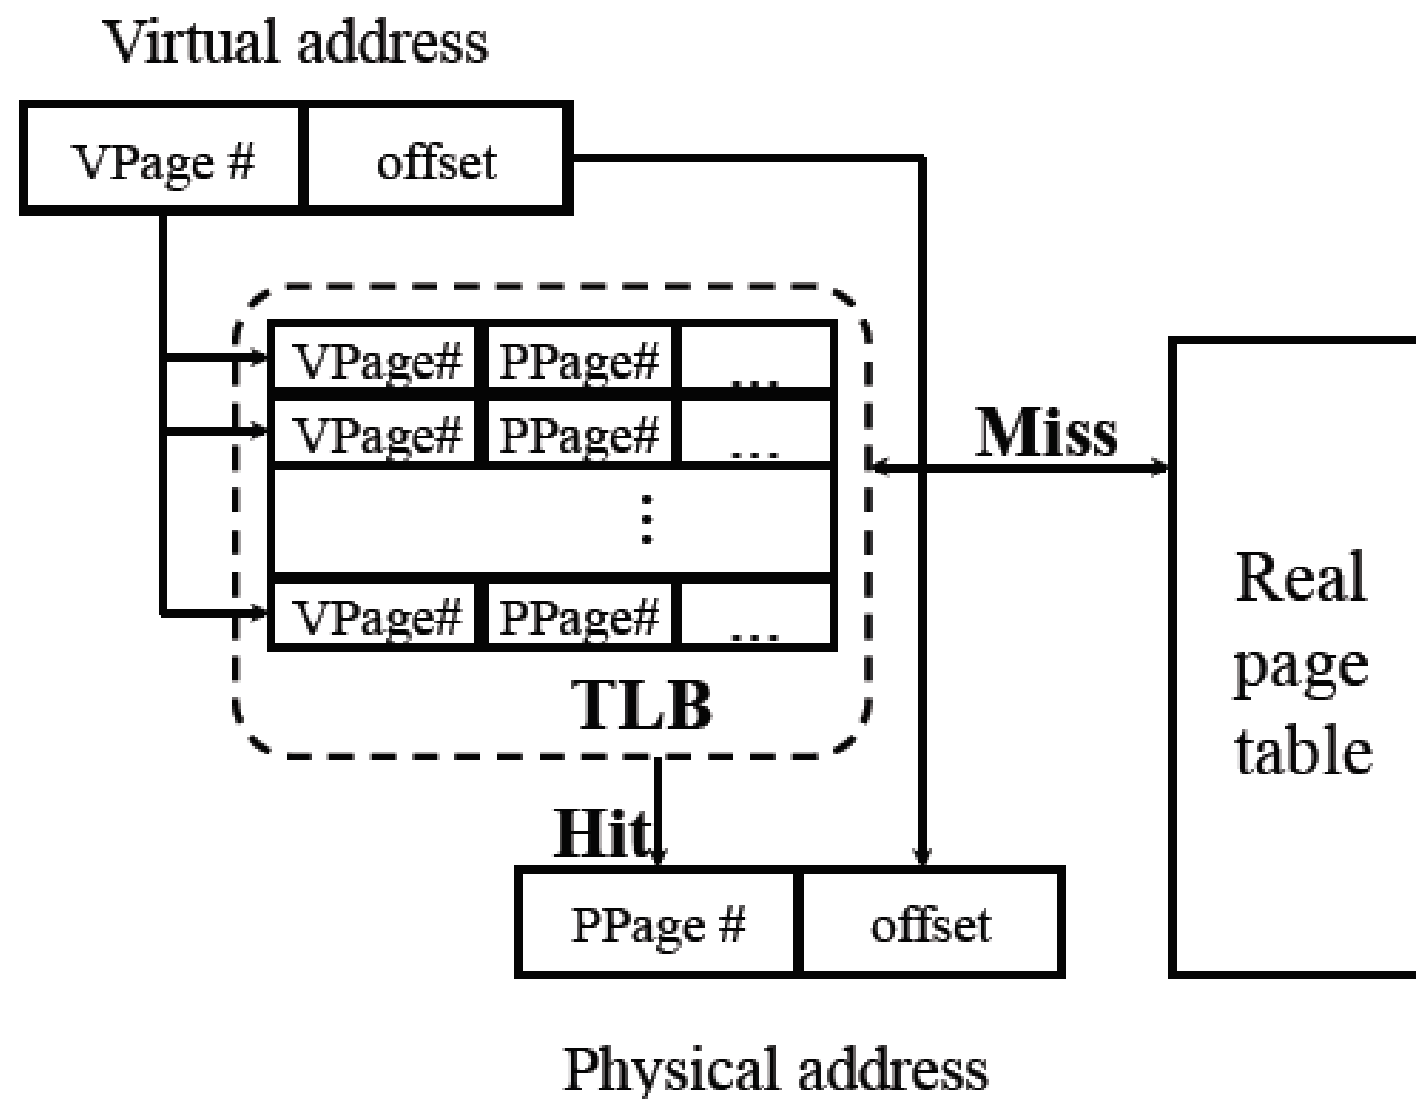


What does this buy us?

Virtual-To-Physical Lookups

- Programs only know virtual addresses
 - Each program or process starts from 0 to high address
- Each virtual address must be translated
 - May involve walking through the hierarchical page table
 - Since the page table stored in memory, a program memory access may requires several actual memory accesses
- Solution
 - Cache “active” part of page table in a very fast memory

Translation Look-aside Buffer (TLB)



Bits in a TLB Entry

- Common (necessary) bits
 - Virtual page number: match with the virtual address
 - Physical page number: translated address
 - Valid
 - Access bits: kernel and user (nil, read, write)
- Optional (useful) bits
 - Process tag
 - Reference
 - Modify
 - Cacheable

Hardware-Controlled TLB

- On a TLB miss
 - Hardware loads the PTE into the TLB
 - Write back and replace an entry if there is no free entry
 - Generate a fault if the page containing the PTE is invalid
 - VM software performs fault handling
 - Restart the CPU
- On a TLB hit, hardware checks the valid bit
 - If valid, pointer to page frame in memory
 - If invalid, the hardware generates a page fault
 - Perform page fault handling
 - Restart the faulting instruction

Software-Controlled TLB

- On a miss in TLB
 - Write back if there is no free entry
 - Check if the page containing the PTE is in memory
 - If not, perform page fault handling
 - Load the PTE into the TLB
 - Restart the faulting instruction
- On a hit in TLB, the hardware checks valid bit
 - If valid, pointer to page frame in memory
 - If invalid, the hardware generates a page fault
 - Perform page fault handling
 - Restart the faulting instruction

Hardware vs. Software Controlled

- Hardware approach
 - Efficient
 - Not flexible
 - Need more space for page table
- Software approach
 - Flexible
 - Software can do mappings by hashing
 - $PP\# \rightarrow (Pid, VP\#)$
 - $(Pid, VP\#) \rightarrow PP\#$
 - Can deal with large virtual address space

TLB Related Issues

- What TLB entry to be replaced?
 - Random
 - Pseudo LRU
- What happens on a context switch?
 - Process tag: change TLB registers and process register
 - No process tag: Invalidate the entire TLB contents
- What happens when changing a page table entry?
 - Change the entry in memory
 - Invalidate the TLB entry

Consistency Issues

- “Snoopy” cache protocols (hardware)
 - Maintain consistency with DRAM, even when DMA happens
- Consistency between DRAM and TLBs (software)
 - You need to flush related TLBs whenever changing a page table entry in memory
- TLB “shoot-down”
 - On multiprocessors, when you modify a page table entry, you need to flush all related TLB entries on all processors

Summary – Part 1

- Virtual Memory
 - Virtualization makes software development easier and enables memory resource utilization better
 - Separate address spaces provide protection and isolate faults
- Address translation
 - Base and bound: very simple but limited
 - Segmentation: useful but complex
- Paging
 - TLB: fast translation for paging
 - VM needs to take care of TLB consistency issues