
File Systems

Inf-2201, University of Tromsø
Loic Guegan (loic.guegan@uit.no)

Based on slides from Lars Ailo Bongo(UiT), Kai Li (Princeton University)

Overview

- ▶ Part I
 - ▶ File system abstractions and operations
 - ▶ Protection
 - ▶ File system structure
 - ▶ Disk allocation and i-nodes
 - ▶ Directory and link implementations
 - ▶ Physical layout for performance
- ▶ Part II
 - ▶ Performance and reliability
 - ▶ File buffer cache
 - ▶ Disk failure and file recovery tools
 - ▶ Consistent updates
 - ▶ Transactions and logging

Why Files?


- ▶ Can't we just use main memory?
- ▶ Can't we use a mechanism like swapping to disk?
- ▶ Need to store large amount of information
- ▶ Need the information to survive process termination
- ▶ Need the information to be shareable by processes

Recall Some High-level Abstractions

- ▶ Processes are an abstraction for processors (CPU)
- ▶ Virtual memory is an abstraction for memory
- ▶ File systems are an abstraction for disks (disk blocks)

File System Layers and Abstractions

- ▶ Network file system maps a network file system protocol to local file systems
 - ▶ NFS, CIFS, DAFS, GFS, HDFS, Dropbox, etc
- ▶ Local file system implements a file system on blocks in volumes
 - ▶ Local disks or network of disks
- ▶ Volume manager maps logical volume to physical disks
 - ▶ Provide logical unit
 - ▶ RAID and reconstruction
- ▶ Disk management manages physical disks
 - ▶ Sometimes part of volume manager
 - ▶ Drivers, scheduling, etc



Volume Manager

- ▶ Group multiple disk partitions into a logical disk volume
 - ▶ No need to deal with physical disk, sector numbers
 - ▶ To read a block: `read(vol#, block#, buf, n);`
- ▶ Volume can include RAID, tolerating disk failures
 - ▶ No need to know about parity disk in RAID-5, for example
 - ▶ No need to know about reconstruction
- ▶ Volume can provide error detections at disk block level
 - ▶ Some products use a checksum block for 8 blocks of data
- ▶ Volume can grow or shrink without affecting existing data
- ▶ Volume can have remote volumes for disaster recovery
- ▶ Remote mirrors can be split or merged for backups



Files vs. Block Storage

File abstraction

- ▶ Byte oriented
- ▶ Named files
- ▶ Users protected from each other
- ▶ Robust to machine failures

Disk abstraction

- ▶ Block oriented
- ▶ Block numbers
- ▶ No protection among users of the system
- ▶ Data might be corrupted if machine crashes

File Structure Possibilities


- ▶ **Byte sequence**

- ▶ Read or write a number of bytes
- ▶ Unstructured or linear
- ▶ Unix, Windows




- ▶ **Record sequence**

- ▶ Fixed or variable length
- ▶ Read or write a number of records
- ▶ Not used: punch card days



- ▶ **Tree**

- ▶ Records with keys
- ▶ Read, insert, delete a record (typically using B-tree, sorted on key)
- ▶ Used in mainframes for commercial data processing



File Types, examples

- ▶ ASCII
- ▶ Binary data
 - ▶ Record
 - ▶ Tree
 - ▶ An Unix executable file
 - ▶ header: magic number, sizes, entry point, flags
 - ▶ text
 - ▶ data
 - ▶ relocation bits
 - ▶ symbol table
- ▶ Devices
- ▶ Everything else in the system



Most common file operations

- ▶ Operations for “sequence of bytes” files
 - ▶ Create: create a mapping from a name to bytes
 - ▶ Delete: delete the mapping
 - ▶ Open: authentication, bring key attributes, disk info into RAM
 - ▶ Close: free up table space, force last block write
 - ▶ Seek: jump to a particular location in a file
 - ▶ Read: read some bytes from a file
 - ▶ Write: write some bytes to a file
 - ▶ Get attributes, Set attributes
- ▶ Implementation goal
 - ▶ Operations should have as few disk accesses as possible and have minimal space overhead


Access Patterns

- ▶ **Sequential (the common pattern)**
 - ▶ File data processed sequentially
 - ▶ Examples
 - ▶ Editor writes out a new file
 - ▶ Compiler reads a file
- ▶ **Random access**
 - ▶ Address a block in file directly without passing through predecessors
 - ▶ Examples:
 - ▶ Data set for demand paging
 - ▶ Databases
- ▶ **Keyed access**
 - ▶ Search for a record with particular values
 - ▶ Usually not provided by today's file systems
 - ▶ Examples
 - ▶ Database search and indexing




File System Components

- ▶ **Naming**
 - ▶ File and directory naming
 - ▶ Local and remote operations
- ▶ **File access**
 - ▶ Implement read/write and other functionalities
- ▶ **Buffer cache**
 - ▶ Reduce client/server disk I/Os
- ▶ **Disk allocation**
 - ▶ File data layout
 - ▶ Mapping files to disk blocks
- ▶ **Management**
 - ▶ Tools for system administrators to manage file systems



Steps to Open a file

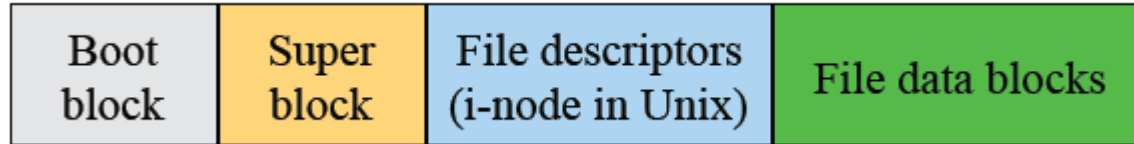
- ▶ File name lookup and authenticate
- ▶ Copy the file descriptors into the in-memory data structure, if it is not in yet
- ▶ Create an entry in the open file table (system wide) if there isn't one
- ▶ Create an entry in PCB
- ▶ Link up the data structures
- ▶ Return a pointer to user



File Read and Write

- ▶ Read 10 bytes from a file starting at byte 2?
 - ▶ seek byte 2
 - ▶ fetch the block
 - ▶ read 10 bytes
- ▶ Write 10 bytes to a file starting at byte 2?
 - ▶ seek byte 2
 - ▶ fetch the block
 - ▶ write 10 bytes in memory
 - ▶ write out the block

Disk Layout




- ▶ **Boot block**
 - ▶ Code to bootstrap the operating system
- ▶ **Super-block defines a file system**
 - ▶ Size of the file system
 - ▶ Size of the file descriptor area
 - ▶ Free list pointer, or pointer to bitmap
 - ▶ Location of the file descriptor of the root directory
 - ▶ Other meta-data such as permission and various times
 - ▶ Kernel keeps in main memory, replicated on disk
- ▶ **File descriptors**
- ▶ **File data blocks**
 - ▶ Data for the files, the largest portion on disk

Data Structures for Disk Allocation

- ▶ The goal is to manage the allocation of a volume
- ▶ A file header for each file
 - ▶ Disk blocks associated with each file
- ▶ A data structure to represent free space on disk
 - ▶ Bit map that uses 1 bit per block (sector)
 - ▶ Linked list that chains free blocks together
 - ▶ ...

111111111111111111110000000000000000
000001111111111100000000000011111111
⋮
1100000111100011110000000000000000



Contiguous Allocation

- ▶ Request in advance for the size of the file
- ▶ Search bit map or linked list to locate a space
- ▶ File header
 - ▶ First block in file
 - ▶ Number of blocks
- ▶ Pros
 - ▶ Fast sequential access
 - ▶ Easy random access
- ▶ Cons
 - ▶ External fragmentation (what if file C needs 3 blocks)
 - ▶ Hard to grow files: may have to move (large) files on disk
 - ▶ May need compaction



Linked Files


- ▶ File header points to 1st block on disk
- ▶ A block points to the next

Pros

- ▶ Can grow files dynamically
- ▶ Free list is similar to a file
- ▶ No external fragmentation or need to move files

Cons

- ▶ Random access: horrible
- ▶ Even sequential access needs one seek per block
- ▶ Unreliable: losing a block means losing the rest



File Allocation Table (FAT)

► Approach


- Table of “next pointers”, indexed by block
 - Instead of pointer stored with block
- Directory entry points to 1st block of file

► Pros

- No need to traverse list to find a block
- Cache FAT table and traverse in memory


► Cons

- FAT table takes lots of space for large disk
 - Hard to fit in memory; so may need seeks
- Pointers for all files on whole disk are interspersed in FAT table
 - Need full table in memory, even for one file
 - Solution: indexed files
 - Keep block lists for different files together, and in different parts of disk




Single-Level Indexed Files

- ▶ A file header holds an array of pointers to point to disk blocks
- ▶ Pros
 - ▶ Can grow up to a limit
 - ▶ Random access is fast
- ▶ Cons
 - ▶ Clumsy to grow beyond the limit
 - ▶ Still lots of seeks



Multi-Level Indexed Files (Unix)

- ▶ 13 Pointers in a header
 - ▶ 1...10: direct pointers
 - ▶ 11: 1-level indirect
 - ▶ 12: 2-level indirect
 - ▶ 13: 3-level indirect
- ▶ Pros & Cons
 - ▶ In favor of small files
 - ▶ Can grow
 - ▶ Limit is 16G and lots of seek




What's in Original Unix i-node?

- ▶ Mode: file type, protection bits, setuid, setgid bits
- ▶ Link count: number of directory entries pointing to this
- ▶ Uid: uid of the file owner
- ▶ Gid: gid of the file owner
- ▶ File size
- ▶ Times (access, modify, change)
- ▶ No filename (why?)
- ▶ 10 pointers to data blocks
- ▶ Single indirect pointer
- ▶ Double indirect pointer
- ▶ Triple indirect pointer

Extents

- ▶ Instead of using a fixed size block, use a number of blocks
 - ▶ XFS uses 8Kbyte block
 - ▶ Max extent size is 2M blocks
- ▶ Index nodes need to have
 - ▶ Block offset
 - ▶ Length
 - ▶ Starting block



Directory Organization Examples

- ▶ **Flat**

- ▶ All files are in one directory

- ▶ **Hierarchical (Unix)**

- ▶ /home/foo/bar
 - ▶ Directory is stored in a file containing (name, i-node) pairs
 - ▶ The name can be either a file or a directory

Mapping File Names to i-nodes

- ▶ **Create/delete**
 - ▶ Create/delete a directory
- ▶ **Open/close**
 - ▶ Open/close a directory for read and write
- ▶ **Link/unlink**
 - ▶ Link/unlink a file
- ▶ **Rename**
 - ▶ Rename the directory



Linear List


- ▶ Method
 - ▶ <FileName, i-node> pairs are linearly stored in a file
 - ▶ Create a file
 - ▶ Append <FileName, i-node>
- ▶ Delete a file
 - ▶ Search for FileName
 - ▶ Remove its pair from the directory
 - ▶ Compact by moving the rest
- ▶ Pros
 - ▶ Space efficient
- ▶ Cons
 - ▶ Linear search
 - ▶ Need to deal with fragmentation

▶ /home/userY/foo/
bar/...
veryLongFileName

```
<foo,1234>  
<bar,1235> ...  
<veryLongFileName,  
4567>
```


Tree Data Structure

- ▶ **Method**
 - ▶ Store <fileName, i-node> a tree data structure such as B-tree
 - ▶ Create/delete/search in the tree data structure
- ▶ **Pros**
 - ▶ Good for a large number of files
- ▶ **Cons**
 - ▶ Inefficient for a small number of files
 - ▶ More space
 - ▶ Complex



Hashing

- ▶ **Method**
 - ▶ Use a hash table to map FileName to i-node
 - ▶ Space for name and metadata is variable sized
 - ▶ Create/delete will trigger space allocation and free
- ▶ **Pros**
 - ▶ Fast searching and relatively simple
- ▶ **Cons**
 - ▶ Not as efficient as trees for very large directory (wasting space for the hash table)



Disk I/Os to Read/Write A File

- ▶ Disk I/Os to access a byte of /home/foo/bar
 - ▶ Read the i-node and first data block of “/”
 - ▶ Read the i-node and first data block of “home”
 - ▶ Read the i-node and first data block of “foo”
 - ▶ Read the i-node and first data block of “bar”
- ▶ Disk I/Os to write a file
 - ▶ Read the i-node of the directory and the directory file.
 - ▶ Read or create the i-node of the file
 - ▶ Read or create the file itself
 - ▶ Write back the directory and the file
- ▶ Too many I/Os to traverse the directory
 - ▶ Solution is to use ***Current Working Directory***



Links

- ▶ **Symbolic (soft) links**


- ▶ A symbolic link is just the name of the file
- ▶ Original owner still owns the file, deleted on rm by owner
- ▶ Use a new i-node for the symbolic link
ln -s source target

- ▶ **Hard links**

- ▶ A link to a file with the same i-node
ln source target
- ▶ Delete may or may not remove the target depending on whether it is the last one (link reference count)

Original Unix File System

- ▶ Simple disk layout
 - ▶ Block size is sector size (512 bytes)
 - ▶ i-nodes are on outermost cylinders
 - ▶ Data blocks are on inner cylinders
 - ▶ Use linked list for free blocks
- ▶ Issues
 - ▶ Index is large
 - ▶ Fixed max number of files
 - ▶ i-nodes far from data blocks
 - ▶ i-nodes for directory not close together
 - ▶ Consecutive blocks can be anywhere
 - ▶ Poor bandwidth (20Kbytes/sec even for sequential access!)




BSD FFS (Fast File System)

- ▶ Use a larger block size: 4KB or 8KB
 - ▶ Allow large blocks to be chopped into fragments
- ▶ Use bitmap instead of a free list
 - ▶ Try to allocate contiguously
 - ▶ 10% reserved disk space



FFS Disk Layout

- ▶ i-nodes are grouped together
 - ▶ A portion of the i-node array on each cylinder
- ▶ Do you ever read i-nodes without reading any file blocks?
- ▶ Overcome rotational delays
 - ▶ Skip sector positioning to avoid the context switch delay
 - ▶ Read ahead: read next block right after the first



What Has FFS Achieved?

- ▶ Performance improvements
 - ▶ 20-40% of disk bandwidth for large files (10-20x original)
 - ▶ Better small file performance
- ▶ We can still do a lot better
 - ▶ Extent based instead of block based
 - ▶ Use a pointer and size for all contiguous blocks (XFS, Veritas file system, etc)
 - ▶ Synchronous metadata writes hurt small file performance
 - ▶ Asynchronous writes with certain ordering (“soft updates”)
 - ▶ Logging (talk about this later)
 - ▶ Play with semantics (/tmp file systems)



Side note : Protection Policy vs. Mechanism

- ▶ A protection system is the mechanism to enforce a security policy
 - ▶ Roughly the same set of choices, no matter what policy
- ▶ A security policy determines what is acceptable or not
 - ▶ Example security policies:
 - ▶ Each user can only allocate 40GB of disk
 - ▶ No one but root can write to the password file
 - ▶ You cannot read my mail

Protection Mechanisms

▶ Authentication

- ▶ Make sure system knows whom it is talking to
 - ▶ Unix: password
 - ▶ US banks: account # + last transactions
 - ▶ Bars: driver's license

▶ Authorization

- ▶ Determine if “X” is allowed to do “Y”
- ▶ Need a simple database

▶ Access enforcement

- ▶ Enforce authorization decision
- ▶ Must make sure there are no loopholes
- ▶ Hard to assert

Protection Domain

- ▶ A set of (objects, rights) pairs
 - ▶ Domain may correspond to single user, or more general
 - ▶ Process runs in a domain at a given instant in time
- ▶ Once identity known, what is Bob allowed to do?
 - ▶ More generally: must be able to determine what each “principal” is allowed to do with what
- ▶ Can be represented as a “protection matrix” with one row per domain, one column per resource
- ▶ What are the pros and cons of this approach?

	File A	Printer B	File C
Domain 1	R	W	RW
Domain 2	RW	W	...
Domain 3	R	...	RW

Access Control Lists (ACLs)

- ▶ By column: For each object, indicate which users are allowed to perform which operations
 - ▶ In most general form, each object has a list of `<user,privileged>` pairs
- ▶ Access control lists are simple, and are used in almost all file systems
 - ▶ Owner, group, world
- ▶ Implementation
 - ▶ Stores ACLs in each file
 - ▶ Use login authentication to identify
 - ▶ Kernel implements ACLs

Capabilities

- ▶ By rows: For each user, indicate which files may be accessed and in what ways
 - ▶ Store a lists of <object, privilege> pairs for each user.
 - ▶ Called a *Capability List*
- ▶ Capabilities frequently do both naming and protection
 - ▶ Can only “see” an object if you have a capability for it.
 - ▶ Default is no access
- ▶ Implementation
 - ▶ Capability lists
 - ▶ Architecture support
 - ▶ Stored in the kernel
 - ▶ Stored in the user space but in encrypted format
 - ▶ Checking is easy: no enumeration

Access Enforcement

- ▶ Use a trusted party to
 - ▶ Enforce access controls
 - ▶ Protect authorization information
- ▶ Kernel is the trusted party
 - ▶ This part of the system can do anything it wants
 - ▶ If it has a bug, the entire system can be destroyed
 - ▶ Want it to be as small & simple as possible
- ▶ Security is only as strong as the weakest link in the protection system

Summary - Part 1


- ▶ Protection
 - ▶ We basically live with access control list
 - ▶ More protection is needed in the future
- ▶ File system structure
 - ▶ Boot block, super block, file metadata, file data
- ▶ File metadata
 - ▶ Consider efficiency, space and fragmentation
- ▶ Directories
 - ▶ Consider the number of files
- ▶ Links
 - ▶ Soft vs. hard
- ▶ Physical layout
 - ▶ Where to put metadata and data

Overview

- ▶ Part I
 - ▶ File system abstractions and operations
 - ▶ Protection
 - ▶ File system structure
 - ▶ Disk allocation and i-nodes
 - ▶ Directory and link implementations
 - ▶ Physical layout for performance
- ▶ Part II
 - ▶ Performance and reliability
 - ▶ File buffer cache
 - ▶ Disk failure and file recovery tools
 - ▶ Consistent updates
 - ▶ Transactions and logging

File Buffer Cache for Performance

- ▶ Cache files in main memory
 - ▶ Check the buffer cache first
 - ▶ Hit will read from or write to the buffer cache
 - ▶ Miss will read from the disk to the buffer cache
- ▶ Usual questions
 - ▶ What to cache?
 - ▶ How to size?
 - ▶ What to prefetch?
 - ▶ How and what to replace?
 - ▶ Which write policies?



What to Cache?

- ▶ Things to consider
 - ▶ I-nodes and indirect blocks of directories
 - ▶ Directory files
 - ▶ I-nodes and indirect blocks of files
 - ▶ Files
- ▶ What is a good strategy?
 - ▶ Cache i-nodes and indirect blocks if they are in use?
 - ▶ Cache only the i-nodes and indirect blocks of the current directory?
 - ▶ Cache an entire file vs. referenced blocks of files?



How to Size?


- ▶ An important issue is how to partition memory between the buffer cache and VM cache
- ▶ Early systems use fixed-size buffer cache
 - ▶ It does not adapt to workloads
- ▶ Later systems use variable size cache
 - ▶ But, large files are common, how do we make adjustment?
- ▶ Basically, we solve the problem using the *working set idea*

What to Prefetch?

- ▶ **Optimal**
 - ▶ The blocks are fetched in just enough time to use them
 - ▶ But, too hard to do
- ▶ **The good news is that files also have locality**
 - ▶ Temporal locality
 - ▶ Spatial locality
- ▶ **Common strategies**
 - ▶ Prefetch next k blocks together (typically $> 64\text{KB}$)
 - ▶ Some discard unreferenced blocks
 - ▶ Cluster blocks of the same directory and i-nodes if possible (to the same cylinder group and neighborhood) to make prefetching efficient


How and What to Replace?

- ▶ Page replacement theory
 - ▶ Use past to predict future
 - ▶ LRU is good
- ▶ Buffer cache with LRU replacement mechanism
 - ▶ If b is in buffer cache, move it to front and return b
 - ▶ Otherwise, replace the tail block, get b from disk, insert b to the front
 - ▶ Use double linked list with a hash table



Which Write Policies?

- ▶ **Write through**
 - ▶ Whenever modify cached block, write block to disk
 - ▶ Cache is always consistent
 - ▶ Simple, but cause more I/Os
- ▶ **Write back**
 - ▶ When modifying a block, mark it as dirty & write to disk later
 - ▶ Fast writes, absorbs writes, and enables batching
 - ▶ So, what's the problem?





Write Back Complications

- ▶ **Fundamental tension**
 - ▶ On crash, all modified data in cache is lost.
 - ▶ The longer you postpone write backs, the faster you are but the worst the damage is on a crash
- ▶ **When to write back**
 - ▶ When a block is evicted
 - ▶ When a file is closed
 - ▶ On an explicit flush
 - ▶ When a time interval elapses (30 seconds in Unix)
- ▶ **Issues**
 - ▶ These write back options have no guarantees
 - ▶ A solution is consistent updates (later)

File Recovery Tools

- ▶ **Physical backup (dump) and recovery**
 - ▶ Dump disk blocks by blocks to a backup system
 - ▶ Backup only changed blocks since the last backup as an incremental
 - ▶ Recovery tool built accordingly
- ▶ **Logical backup (dump) and recovery**
 - ▶ Traverse the logical structure from the root
 - ▶ Selectively dump what you want to backup
 - ▶ Verify logical structures as you backup
 - ▶ Recovery tool selectively move files back
- ▶ **Consistency check (e.g. fsck)**
 - ▶ Start from the root i-node
 - ▶ Traverse the whole tree and mark reachable files
 - ▶ Verify the logical structure
 - ▶ Figure out what blocks are free




What fsck does

- ▶ Get default list of file systems to check from */etc/fstab*
- ▶ Inconsistencies checked:
 - ▶ Blocks claimed by more than one i-node or the free map
 - ▶ Blocks claimed by an i-node outside range of the filesystem
 - ▶ Incorrect link counts
 - ▶ Size checks (directory size etc)
 - ▶ Bad i-node format
 - ▶ Blocks not accounted anywhere
 - ▶ Directory checks:
 - ▶ File pointing to unallocated i-node; I-node number out of range; . or .. Not first two entries of a directory or have wrong i-node number
 - ▶ Super Block checks
 - ▶ More blocks for i-nodes than are in the filesystem; Bad free block map format; Total free block and/or free i-node count incorrect
 - ▶ Put orphaned files and directories in lost+found directory



Persistence and Crashes

- ▶ File system promise:
Persistence
 - ▶ File system will hold a file until its owner explicitly deletes it
- ▶ Why is this hard?
 - ▶ A crash will destroy memory content
 - ▶ Cache more \Rightarrow better performance
 - ▶ Cache more \Rightarrow lose more on a crash
 - ▶ A file operation often requires modifying multiple blocks, but the system can only atomically modify one at a time
 - ▶ Systems can crash anytime



What is a Crash?


- ▶ Crash is like a context switch
 - ▶ Think about a file system as a thread before the context switch and another after the context switch
 - ▶ Two threads read or write same shared state?
- ▶ Crash is like time travel
 - ▶ Current volatile state lost; suddenly go back to old state
 - ▶ Example: move a file
 - ▶ Place it in a directory
 - ▶ Delete it from old
 - ▶ Crash happens and both directories have problems



Consistent Updates: Problem

◆ Modify /u/cos318/foo

- Traverse to /u/cos318/
Crash → **Consistent**
- Allocate data block
Crash → **Consistent**
- Write pointer into i-node
Crash → **Inconsistent**
- Write new data to foo
Crash → **Consistent**



Writing metadata first can cause inconsistency

Consistent Updates: Data Before Metadata

◆ Modify /u/cos318/foo

- Traverse to /u/cos318/

Crash → **Consistent**

- Allocate data block


Crash → **Consistent**

- Write new data to foo

Crash → **Consistent**

- Write pointer into i-node

Crash → **Consistent**



Consistent Updates: Bottom-Up Order

- ▶ The general approach is to use a “bottom up” order
 - ▶ File data blocks, file i-node, directory file, directory i-node, ...
- ▶ What about file buffer cache?
 - ▶ Write back all data blocks
 - ▶ Update file i-node and write it to disk
 - ▶ Update directory file and write it to disk
 - ▶ Update directory i-node and write it to disk (if necessary)
 - ▶ Continue until no directory update exists
- ▶ Does this solve the write back problem?
 - ▶ Updates are consistent but leave garbage blocks around
 - ▶ May need to run fsck to clean up once a while
 - ▶ Ideal approach: consistent update without leaving garbage



Operations as transactions in FileSys


- ▶ Make a file operation a transaction
 - ▶ Create a file
 - ▶ Move a file
 - ▶ Write a chunk of data
 - ▶ ...
- ▶ Make arbitrary number of file operations a transaction
 - ▶ Just keep logging but make sure that things are idempotent: making a very long transaction
 - ▶ Recovery by replaying the log and correct the file system
 - ▶ This is called logging file system or journaling file system
 - ▶ Almost all new file systems are journaling (Windows NTFS, Veritas file system, file systems on Linux)

Log Management

- ▶ How big is the log? Same size as the file system?
- ▶ Observation
 - ▶ Log what's needed for crash recovery
- ▶ Management method
 - ▶ Checkpoint operation: flush the buffer cache to disk
 - ▶ After a checkpoint, we can truncate log and start again
 - ▶ Log needs to be big enough to hold changes in memory
- ▶ Some logging file systems log only metadata (file descriptors and directories) and not file data to keep log size down

Log-structured File System (LFS)

- ▶ Structure the entire file system as a log with segments
- ▶ A segment has i-nodes, indirect blocks, and data blocks
- ▶ All writes are sequential (no seeks)
- ▶ There will be holes when deleting files



Summary – Part 2

- ▶ File buffer cache
 - ▶ True LRU is possible
 - ▶ Simple write back is vulnerable to crashes
- ▶ Disk block failures and file system recovery tools
 - ▶ Individual recovery tools
 - ▶ Top down traversal tools
- ▶ Logging file systems